Tracing AI Assistance and AI Agent Activity in Survey Research

Valentina González-Rostani* Shir Raviv†

April 8, 2026

Abstract

Generative AI creates an authentication problem for survey research. When respondents can outsource summarization, explanation, or recall to chatbots, an open-ended answer may no longer reflect the respondent's own reasoning. We present an auditable toolkit to detect likely AI-assisted completion by combining response-process paradata with prompt-specific semantic benchmarks. We validate the approach in a randomized survey experiment in which participants complete an article summary task under blocked AI access or observed AI access, and in a second benchmark based on end to end completions by a synthetic AI agent. Across both settings, AI assistance leaves a distinct signature. Responses show thinner traces of drafting and revision, greater similarity to prompt specific AI output, and lower distinctiveness relative to peer responses. These patterns strengthen further for fully automated agent completions. Classifiers built from these signals achieve strong out-of-sample performance.

Keywords: survey methodology; AI; survey authenticity; paradata; embeddings

*University of Southern California, gonzalez.rostani@usc.edu.
†Tel Aviv University.

1 Introduction

Online surveys remain central to empirical research in the social sciences because they allow researchers to recruit large and diverse samples, field complex experiments at relatively low cost, and measure attitudes and preferences at scale (Heath, Fisher, and Smith 2005; Cohen 2020; Thomas 2024). These strengths, however, depend on a basic assumption: recorded answers reflect respondents' own attention, comprehension, and judgment. Generative AI weakens that assumption. Large language models can summarize text, generate explanations, draft open-ended answers, and even complete survey instruments from beginning to end. As a result, a response may look attentive, articulate, and compliant while capturing far less of the respondent's own reasoning than the researcher intends to measure.

Existing safeguards are not well-suited to this problem. Attention checks, CAPTCHAs, and speed flags were built to detect inattentive respondents and simple bots, not respondents who selectively outsource parts of the task to AI or automated agents that can mimic normal survey behavior (Park et al. 2025; Höhne et al. 2025; Westwood, n.d.). Off-the-shelf AI-text detectors offer limited help. Their performance is unstable across prompts, topics, and model versions, and they say little about how a suspicious response was actually produced (Crossley et al. 2024; Yan et al., n.d.; Claassen et al. 2026). The problem is also substantively important rather than hypothetical. Roughly one third of crowdworkers report using LLMs on open-ended survey items (Zhang, Xu, and Alvero 2025); recent work estimates likely AI-assisted responding in live surveys at rates between 4.4% and 45% (Westwood and Frederick, n.d.; Panizza, Kyrychenko, and Roozenbeek 2026); and autonomous AI agents now pass standard survey quality checks at high rates (Westwood, n.d.). Once AI can generate fluent, context-specific text that is difficult to distinguish from human writing (Dathathri et al. 2024), apparent response quality can no longer be treated as evidence of authorship.

This paper addresses that measurement problem with a survey-embedded, auditable

toolkit for tracing AI involvement in open-ended responses. Rather than relying on a generic detector applied after fieldwork, we instrument the response process itself. Our framework combines two layers of evidence. The first is a behavioral layer that captures how answers are produced, including keystrokes, pauses, deletions, revisions, and task timing. The second is a linguistic layer that evaluates the submitted text along four dimensions: AI signature, Evaluation, Creativity and originality, and broader semantic embeddings. AI signature captures how close a response is to benchmark AI output for the same prompt. Evaluation captures features such as typos and related quality assessments produced through LLM-based annotation. Creativity and originality capture how distinctive a response remains relative to peer responses, including measures such as peer distance. Finally, response embeddings capture broader semantic properties of the text itself. Open-ended items are especially useful in this setting because they preserve traces of drafting and revision while also generating text that can be evaluated for similarity, originality, and source alignment.

Our empirical strategy rests on three design choices. First, we study an open-ended article-summary task that is common in survey experiments and well suited to measuring both final text and response process. Participants recruited through CloudResearch Connect read one of seven short articles and then wrote a brief summary and title in their own words. Second, we randomize AI access and implement it through a dedicated assistant embedded inside the survey. The tool was built in Streamlit and connected to OpenAI's `gpt-4o-mini` model, which allows us to observe actual AI use rather than infer it after the fact. Third, we administer the same instrument to a synthetic browser-based AI agent and collect over 500 end-to-end agent completions. Together, these design features let us compare unaided human responses, AI-assisted human responses, and fully automated responses within a common survey environment.

Our findings show that AI assistance changes both what respondents submit and how they produce it. Relative to the blocked-AI control group, AI-assisted respondents submit longer and cleaner answers, but those answers show much thinner traces of drafting

and revision: paste activity increases, keystroke-to-character ratios and deletion rates fall, responses move closer to prompt-specific AI benchmarks, and responses become less distinctive relative to peers. Fully automated agent responses intensify the same pattern, showing the weakest production traces, the greatest similarity to benchmark AI output, and the lowest peer originality. When we combine these behavioral and semantic measures in machine-learning classifiers, out-of-sample performance is strong. The best models exceed 0.9 ROC-AUC, and the top random-forest specification reaches 0.919 ROC-AUC and 0.971 average precision when distinguishing AI-assisted or agent responses from the control group. Even without considering behavioral indicators, which may become less useful as agents learn to better mimic human response behavior, AI use remains detectable through text embeddings that capture similarity to generic benchmark AI output and creativity relative to peer responses.

This paper makes two contributions. Empirically, it shows that human, human-AI, and fully automated survey responses leave systematically different behavioral and textual traces, even when final answers appear fluent and on topic. Methodologically, it offers a transparent workflow that researchers can deploy in online surveys by combining front-end instrumentation, response-process paradata, prompt-specific semantic benchmarks, and classification models. As a practical contribution, we also provide predictive models that scholars can use to distinguish human, AI-assisted, and fully AI-generated responses. Finally, we show how synthetic AI agents can be used before fielding as a stress test for survey instruments, helping researchers identify prompts that are especially easy to outsource. More broadly, we argue that, in the age of generative AI, authorship is not a background assumption but part of the measurement problem itself.

2 A Toolkit for Tracing AI-Assisted Responses

How should researchers study and address AI assistance in surveys intended to measure human responses? This question calls for a broader approach than a generic text detec-

tor or a single threshold, because contamination can take several forms: direct copying, partial borrowing, scaffolded drafting, or heavily edited transcription. The recent literature has made important progress in identifying AI-generated and AI-assisted survey responses, but most existing approaches still rely on one of two partial strategies. One focuses on the text itself, using classifiers to distinguish human from machine-written responses based on linguistic or semantic features (Claassen et al. 2026; Kosmyna et al. 2025; Zhang, Xu, and Alvero 2025). These methods can perform well within specific prompts or corpora, but their accuracy often declines across settings, and they reveal little about how a response was produced. The other focuses on a narrow set of behavioral signals, such as paste events or unusually rapid typing (Westwood and Frederick, n.d.; Deane et al. 2026; Diedenhofen and Musch 2017). These indicators can provide useful initial evidence, but on their own they may miss more selective forms of assistance, including cases in which respondents consult AI and then retype or lightly edit the output before submitting it (Westwood and Frederick, n.d.).[1]

Our approach builds on these earlier efforts by combining multiple sources of evidence within a single, auditable framework. Specifically, we integrate preventive design, process paradata, and task-specific semantic diagnostics into a toolkit that helps researchers both identify likely AI assistance transparently and diagnose where the survey itself may be inviting delegation. We organize the toolkit around four core domains: *Behavior*, *AI signature*, *Evaluation*, and *Creativity and originality*. We then extend this core in two ways. First, we incorporate *Text embeddings*, which capture broader semantic properties of the response itself. Second, in tasks where respondents write about a shared source text, we add *Text alignment*, which measures how closely the response tracks the assigned material. The aim is not to produce a universal black-box verdict, but to preserve a transparent set of indicators that allows researchers to explain why a response appears authentic, suspicious, or difficult to classify.

Open-ended questions are especially useful for this purpose because they preserve

1. For an extended discussion of existing approaches, refer to Appendix A.1.

evidence not only in the final response but also in the process that produced it. Unlike closed-ended items, they allow researchers to observe traces of composition, revision, and transfer while also evaluating the content of the submitted text. The resulting evidence is therefore inherently multidimensional. It includes how the response was written, how closely it resembles a prompt-specific AI benchmark, how distinctive it is relative to peer responses, whether it appears unusually polished, and, when respondents write about a shared source, how closely it remains anchored to the assigned text. Table 1 summarizes these families of measures and their role in the workflow.[2]

Table 1: Toolkit measure families

Measure family	Illustrative measures	Authentication role
Behavior	Paste indicator, paste count, task time, seconds per word, keystrokes, extra keys, deletions, key-to-character ratio	Captures whether the response appears to have been drafted inside the survey environment rather than transferred in or produced through unusually linear composition.
Text embedding	Summary-text embedding dimensions, sentence embedding representation of the response	Encodes the response itself as a dense semantic representation, allowing the classifier to detect broader semantic regularities not captured by targeted hand-crafted measures.
AI signature	Similarity to benchmark AI answer, title similarity to benchmark AI answer, novelty relative to AI benchmark	Measures proximity to a prompt-specific machine counterfactual rather than to an abstract notion of "AI writing."
Evaluation	Quality index, clarity/coherence, accuracy, typo indicators	Helps interpret whether quality gains are accompanied by other signs of outsourcing; used as supplementary evidence rather than a standalone detector.
Creativity and originality	Peer originality, peer distance, creativity index, lexical diversity, n-gram novelty	Captures whether a response remains distinctive relative to other respondents or collapses toward highly standardized output.
Text alignment	Similarity to text reference (if available)	Evaluates whether the response remains anchored to assigned source material. This category is task-specific and is only applicable when respondents are writing about a common reference text.

Behavior. The first category captures how the response was produced. It includes direct indicators of transfer, such as paste events, along with broader measures of writing

2. Appendix A.3 contains further details on the measurement pipeline.

behavior, including task time, seconds per word, keystrokes, extra keystrokes, deletion activity, and the keystroke-to-character ratio. No single indicator is sufficient on its own. Together, however, these measures help distinguish responses that were drafted inside the survey box through a process of writing, revision, and correction from responses that appear to have been transferred or transcribed with relatively little visible composition. These measures are especially useful because they are easy to interpret and map closely onto the core question of interest: whether the respondent actually composed the answer within the survey environment.

Text embeddings. The second category captures the overall semantic profile of the response. In practice, we convert each summary into a numerical representation of its meaning and use that representation to detect broader patterns in the text. This is useful because respondents may revise AI-generated text enough that it no longer appears copied or closely matches a benchmark answer, while still retaining broader features of AI-assisted writing.[3,4] The advantage of this measure is that it captures information that simpler indicators may miss.

Evaluation. The *Evaluation* category captures the apparent quality of the final response, including clarity, coherence, accuracy, and typos. These measures matter because AI assistance may improve the submitted answer even when the writing process itself shows limited evidence of independent composition. Their role is not to treat strong writing as suspicious on its own, but to identify cases in which the final response appears unusually polished relative to the effort visible in the behavioral trace. In that sense, evaluation measures help describe the kind of quality gains associated with AI assistance. We

3. The use of embeddings is increasing in political science. It has been used to recover latent political concepts, place actors in a common semantic space, and support inference about substantively meaningful semantic differences across groups and periods (Rheault and Cochrane 2020; Rodriguez, Spirling, and Stewart 2023; Lin 2025).

4. We encode each respondent's text using the sentence-transformer model `sentence-transformers/all-MiniLM-L6-v2`, which produces a 384-dimensional embedding. Inputs are truncated at 256 tokens, pooled using attention-mask mean pooling over token representations, and L2-normalized before entering the analysis.

treat them as supplementary evidence rather than standalone detectors, and we estimate them using an LLM judge through the OpenAI API.

AI signature. The *AI signature* category captures how closely a respondent's answer resembles a benchmark AI response generated for the same prompt. For each task, we generate an AI answer using the same article and instructions shown to respondents, and then compare each submitted response to that benchmark in a shared semantic space through comparing text embeddings (see Figure 1). The resulting measures indicate whether a summary or title falls unusually close to the kind of answer an AI model would produce under those exact conditions. This matters because some prompts tightly constrain the space of reasonable responses. A generic detector cannot distinguish between answers that appear machine-like because they were AI-assisted and answers that look similar simply because the task leaves limited room for variation. By grounding the comparison in the specific prompt, article, and instructions, this category provides a transparent and task-specific measure of AI resemblance.

Creativity and originality. The *Creativity and originality* captures whether a response remains distinctive relative to other answers to the same prompt. It includes both semantic and surface-level measures. On the semantic side, *peer originality* compares each response to the center of the remaining responses for that prompt (embeddings of each respondent relative to the centroid of embeddings from other participants), so that higher values indicate greater distinctiveness from the local distribution. Figure 1 clarifies the logic of the two main embedding-based comparisons in the toolkit: *AI signature* and *peer originality*. Both begin by mapping a respondent's answer into a common embedding space. *AI signature* compares that answer to a benchmark AI response generated for the same prompt, producing a task-specific measure of proximity to AI output. *Peer originality*, by contrast, compares the same answer to the center of peer responses to the prompt, producing a measure of how unusual it is relative to the local response distribution. Taken together, these two measures distinguish closeness to a benchmark AI answer

from distinctiveness relative to other human responses.

To complement this measure, we incorporate lexical diversity measures such as the type-token ratio and Yule's K, which capture how repetitive or varied the vocabulary is, while novelty measures based on uncommon n-grams capture whether the phrasing is less formulaic relative to the sample. These measures are useful because AI assistance does not always appear as direct copying or obvious transfer. Respondents may borrow, scaffold, or lightly revise AI-generated text in ways that make responses more standardized without leaving a clear paste signature. This category, therefore, asks whether a response remains semantically and linguistically distinctive.

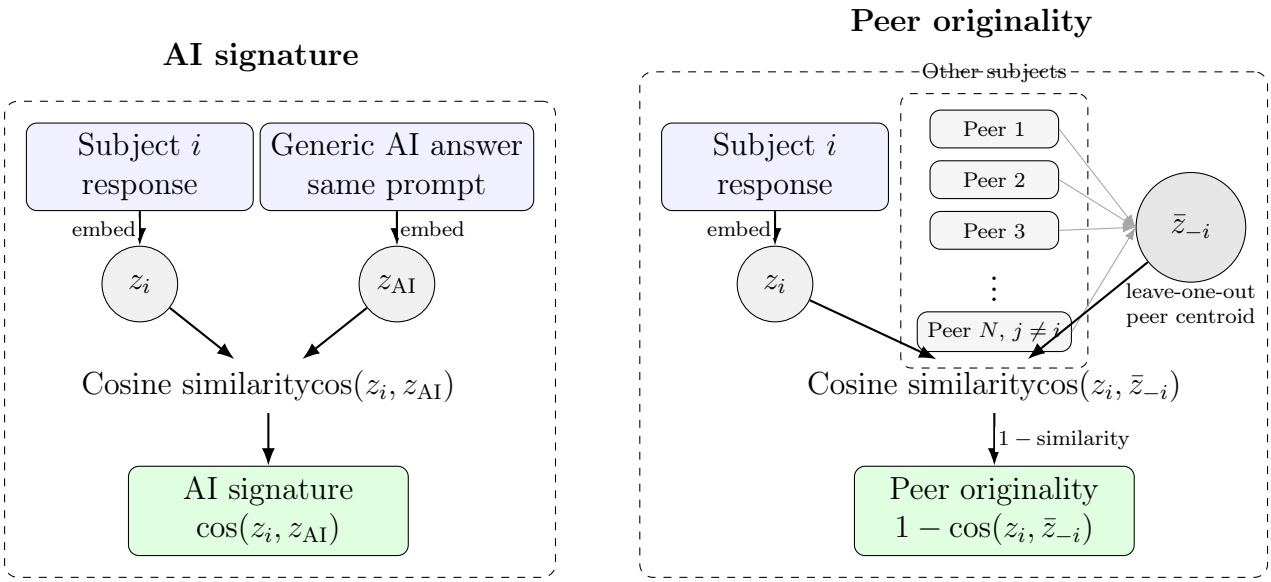

Figure 1: Construction of prompt-specific semantic measures for AI proximity and respondent distinctiveness.

Notes: Both measures begin by mapping a subject's response into a common embedding space. AI signature compares that embedding to the embedding of a generic AI answer generated under the same prompt. Peer originality compares the same subject embedding to the leave-one-out centroid of all other respondents' embeddings and is defined as one minus cosine similarity, so that larger values indicate greater distinctiveness relative to peers.

Text alignment. The final category, *Text alignment*, is task-specific by design. In our setting, respondents are asked to summarize a common source text, so it is substantively meaningful to compare each response to the article itself. Measures such as similarity to the article and faithfulness to the source capture whether the response remains anchored to the assigned text rather than drifting toward generic or externally generated content.

In many open-ended survey settings, there is no common reference text, and in those cases, text alignment is not applicable. The goal is not to impose a single template on all open-ended items, but to incorporate source-based diagnostics when the design makes them possible.

Taken together, these categories produce an auditable profile rather than a black-box judgment. A respondent may use AI and then retype or lightly revise the output, reducing obvious signs of transfer while still leaving a response that is unusually linear in its production and unusually close to a benchmark AI answer. Another respondent may submit a polished and accurate response with strong evaluator scores, yet still show a broader pattern consistent with AI assistance. In tasks built around a shared source text, responses may also stray from the assigned material even when they remain fluent and superficially plausible. The strength of the toolkit therefore lies in combination. No single indicator is decisive on its own. What matters is the cumulative pattern across multiple, interpretable forms of evidence tied to how open-ended survey responses are actually produced.

3 Validation

We validate the toolkit in a randomized survey experiment that introduces AI access under controlled conditions and then tests whether the indicators in our framework capture the resulting differences in response production. We recruited participants through CloudResearch Connect, a widely used online platform for survey experiments in political behavior (e.g., Stagnaro et al. 2024). Participants completed a short survey that included an open-ended task: they read a brief article and then wrote a short summary in their own words.

We chose this task because it closely matches common practice in political behavior research. Many survey experiments expose respondents to a short text and then use follow-up items to assess whether they attended to the information, understood it, and

could restate its main message (e.g., Ziegler 2022; Kane and Barabas 2019). The value of such a task rests on a key assumption: that respondents actually processed the article and produced the response on their own. As AI tools make summarization easy to outsource, that assumption becomes less secure. A response may appear fluent and accurate while revealing less about the respondent's own effort and comprehension. The summary task therefore provides a realistic and substantively meaningful setting for studying AI-assisted responding.

We randomly varied whether respondents were given access to AI while completing the task. In the control condition, respondents were instructed not to use external AI tools. In the treatment condition, respondents were informed that an AI writing assistant was available and that its use was entirely optional. To create a controlled setting in which AI access was observed rather than inferred post hoc, we developed our own LLM-based app and embedded it directly in the survey using Streamlit and the OpenAI API. This allowed us to track tool use and link it to individual respondents. By embedding the tool within the survey, we can more directly assess how access to AI affects both the response process and the text respondents ultimately submit.[5]

We also randomized the topic of the article to be summarized, assigning each respondent to one of seven short texts. This reduces the risk that the results are driven by idiosyncratic features of any single article and helps ensure that the patterns we identify are not specific to one topic or prompt. Appendix A.2 provides additional details on the experimental design and reproduces the full text of all task materials.

To operationalize the toolkit, we instrument the open-ended task to capture both process- and text-based outcomes. On the process side, we record paste attempts, deletion and revision behavior, keystroke activity, and response timing. On the text side, we

5. As part of a separate preregistered project on how AI use shapes policy attitudes, the broader experiment also included two additional treatment arms that varied the framing and perceived quality of the AI tool. Because those conditions served different substantive purposes, we focus here on the cleanest comparison: respondents assigned to the blocked-AI control condition (n=412) and respondents assigned to the standard AI-access condition (n=416). Appendix A.7 reports results for the full sample (n=1,647), controlling for the additional treatment arms, and shows that the main patterns are unchanged in both direction and magnitude.

analyze the summary using measures of similarity to benchmark AI outputs, to the source article, and to peer responses, respectively. Our main validation test examines whether experimentally introduced AI access alters how respondents produce open-ended responses. In later sections, we extend this analysis by asking whether the same indicators also distinguish fully automated survey completion by an AI agent.

3.1 Behavioral and Textual Effects of AI Access

Figure 2 displays the effects of AI assistance relative to the blocked-AI control group, expressed as topic-adjusted standardized contrasts. Across these outcomes, AI access is associated with longer submitted text, substantially more paste activity, fewer keystrokes relative to the amount of text retained, and shorter completion times. The magnitude of these differences is also large. For example, AI assistance reduces the keystrokes-per-character ratio by 1.040, a drop more than twice the control-group level. This indicates a strong shift away from in-box drafting and toward paste-based production, where much of the final text is inserted rather than typed. The key point is not simply that assisted responses look different at the point of submission. They are produced through a different workflow. Relative to blocked access, the behavioral trace shifts away from iterative drafting inside the survey box and toward transfer, light editing, and efficient transcription of externally generated text.

The semantic plots point in the same direction. Figure 2 shows that AI-exposed responses move closer to the study-specific AI benchmark and that peer distinctiveness also decline. Furthermore, the evaluator identifies fewer typos in AI-assisted answers relative to the control. This is precisely the pattern that creates a survey-validity problem: the final answer can remain on topic and look competent while becoming less diagnostic of the respondent's own processing.

AI access often makes responses more homogeneous, even though the language can be more diverse; they are well-drafted, longer, and reduce evidence of respondent effort. If a researcher looked only at the final text, AI-exposed respondents could appear more

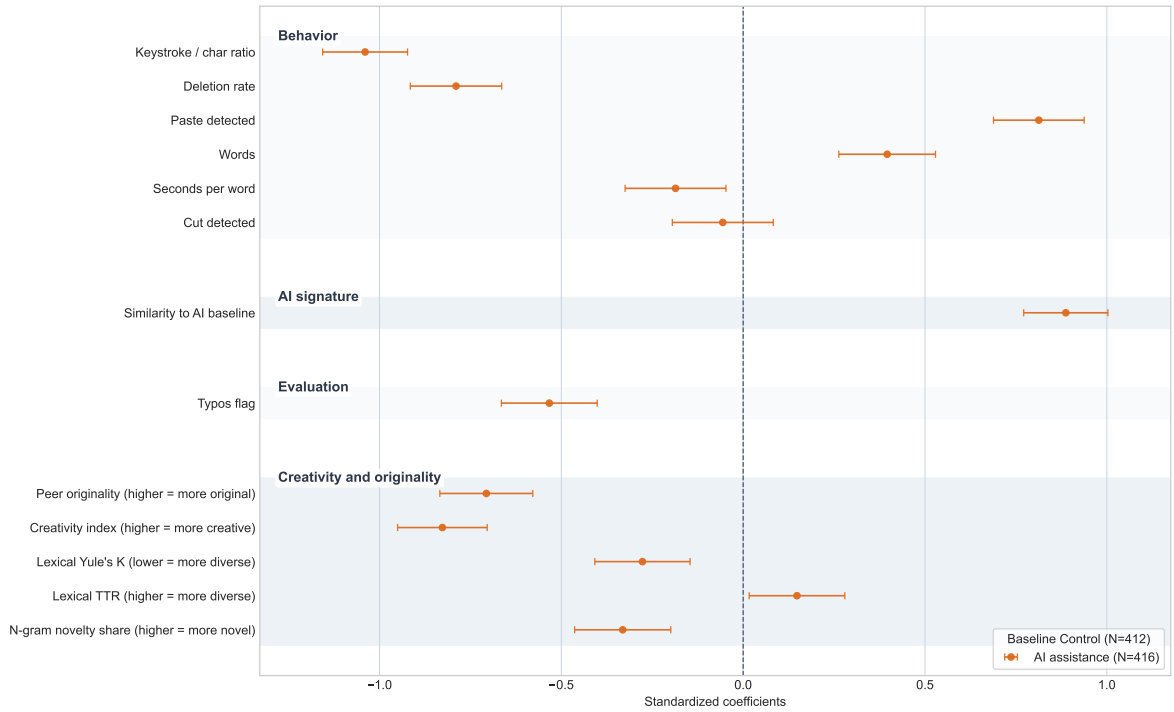

Figure 2: Standardized differences in summary-task outcomes between the human AI-only treatment and the human control group.

Notes: Each point is the coefficient on an indicator for assignment to the AI assistance arm from a separate topic-adjusted OLS regression with article-topic fixed effects; horizontal bars show 95% HC1 confidence intervals. All dependent variables are standardized, so coefficients are in standard-deviation units rather than raw units. Positive coefficients indicate higher values in `AI assisted` than in the baseline control group; negative coefficients indicate lower values. The *Behavior* category captures the writing process: *Words* is final summary length, *Seconds per word* captures production speed, *Keystroke / char ratio* and *Deletion rate* capture editing effort, and *Paste detected* and *Cut detected* are binary indicators for whether those actions occurred during writing. The *AI signature* measure is similarity to an AI baseline summary. The *Evaluation* measure, *Typos flag*, is a binary indicator coded one when evaluators flagged typos. The *Creativity and originality* outcomes capture peer-rated originality and creativity, lexical diversity (*Lexical TTR*; *Yule's K*, where lower raw values imply greater diversity), and n-gram novelty. Refer to Figure A.5 for additional measures, and refer to Figure A.4 for other AI arms, such as low-quality AI provided. Refer to the full table in Table A.1.

articulate, more accurate, or more engaged. The paradata and novelty suggest a different interpretation: some of the apparent improvement reflects changes in production conditions rather than deeper processing of the stimulus.

3.2 Validation with a Synthetic AI Agent

To extend the validation beyond human respondents who could use AI tools, we administered the same survey to a `SyntheticSurveyAgent` benchmark respondent. Previous work has shown that LLMs can generate synthetic respondents, and we build on that line of research by using an agent to probe our survey design (Argyle et al. 2023; Suguri Motoki et al. 2024; Sarstedt et al. 2024). Our objective, however, is not population

emulation. Instead, we incorporate the agent into the design stage of the survey itself. Related studies have used LLMs to evaluate the quality of survey items and responses (Fuchs, Haensch, and Weber 2026; Terry et al. 2025). We instead use AI agents to study how behavior and outputs change across different conditions of AI assistance.

The agent was run through the Chrome browser channel and used `gpt-4o-mini` to answer survey items and generate open-text responses. Each synthetic completion was anchored to a pre-specified persona profile so that the benchmark would reflect a plausible survey participant rather than an unconditioned model producing abstract text. We obtained 543 completed agent responses. For benchmarking, we compare those responses with the same human control group used as the non-AI baseline and with the human AI-assistance group. This puts fully automated completion and assisted human completion on the same scale and shows whether the toolkit is recovering a common response signature rather than overfitting to a single mode of AI use.

Figure 3 shows a consistent ordering across the most diagnostic measures. Human AI assistance already lowers keystroke-to-character ratios, reduces deletion, increases paste incidence, and raises similarity to the AI benchmark relative to unaided human writing. The AI agent intensifies those same patterns: it leaves the thinnest drafting trace, shows the greatest benchmark similarity, and exhibits the lowest distinctiveness relative to peer responses. In many behavioral measures, the agent's effect relative to the control group is roughly twice as large as the effect of human AI assistance. By contrast, indicators based on similarity to the AI benchmark show little difference between the AI-assisted human group and the AI agent, which is consistent with the fact that both rely on AI-generated text. Measures based on peer distance and creativity, however, separate the groups more clearly, with the AI agent appearing even less distinctive relative to peer responses than AI-assisted writing.

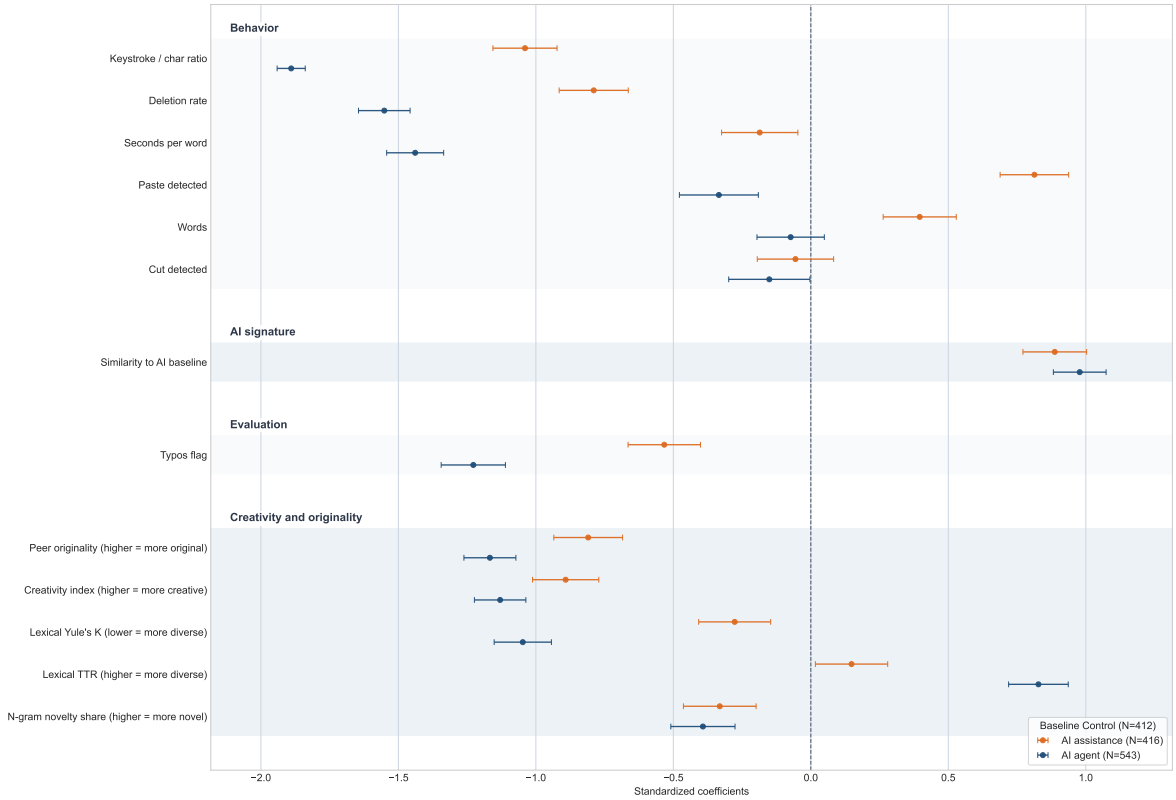

Figure 3: Standardized differences in summary-task outcomes for AI-assisted human responses and AI-agent responses, each relative to the human control group.
Notes: The figure overlays two separate contrasts against the same baseline human control group. The *AI assistance* series compares human responses in the AI-only treatment to the control group. The *AI agent* series compares agent-generated responses to the same control group. Each point is estimated from a separate topic-adjusted OLS model with article-topic fixed effects, and horizontal bars indicate 95% HC1 confidence intervals. All dependent variables are standardized within the relevant comparison sample, so coefficients can be interpreted as differences in standard-deviation units. Positive coefficients indicate higher values than the human control group within that comparison, whereas negative coefficients indicate lower values. The outcome categories are the same as in Figure 2: *Behavior* captures writing effort, speed, and editing actions; *AI signature* captures similarity to an AI baseline summary; *Evaluation* includes the typo-flag indicator; and *Creativity and originality* captures originality, creativity, lexical diversity, and novelty. Refer to the full table in Table A.2.

4 Classification and Predictive Performance

Besides introducing the toolkit, we also train a model to predict AI usage. Researchers could use the toolkit on its own, but the machine learning exercise helps identify which indicators are most informative for detection. The best out-of-sample models exceed 0.9 ROC-AUC (see Figure 4). Here, the top-performing model is a random forest, which combines evaluation, AI-signature measures (similarity to generic AI output), and behavioral indicators. The strongest standalone signals are behavioral, followed by AI signature and

peer originality.[6] We also implement a sensitivity test that excludes behavioral indicators and instead aims to distinguish AI agents from all human-collected data, pooling unassisted and AI-assisted human responses into a single comparison group.[7] Even under this restriction, we are able to train a high-accuracy classifier, with the best-performing model relying on AI signature, evaluation, and creativity indicators (see Figure A.10).

The appropriate practical use is not to impose a universal exclusion threshold. A better approach is to calibrate study-specific cutoffs on transparent indicators or to generate an AI-likely score from those same measures and then report how the rule was chosen. The workflow is most defensible when the analyst can show, before seeing the substantive results, which indicators matter, why they matter, and how the rule performs in validated data. That approach fits contemporary standards for survey-quality reporting better than reliance on a proprietary detector with no clear decision logic.

5 Implications for Survey Research

One of the central challenges now facing online survey research is establishing whether responses reflect respondents' own effort and judgment. As generative AI makes summarization, explanation, and drafting increasingly easy to outsource, survey quality becomes in part an authentication problem. This paper provides steps toward addressing that problem by embedding authentication directly into the survey itself. By combining front-end instrumentation, writing-process paradata, prompt-specific AI benchmarks, and classification models, we show how researchers can recover evidence not only about what respondents submit, but also about how those responses are produced.

We find that AI-assisted and fully automated responses leave identifiable traces. Responses produced with AI are often longer, cleaner, and closer to benchmark AI output, yet they also show thinner traces of drafting and revision and lower distinctiveness rela-

6. Appendix Table A.3 reports the standalone contribution of each indicator.

7. This sensitivity analysis is important as individuals creating AI agents may become more experts on its design and could eventually mimic human behavior when writing.

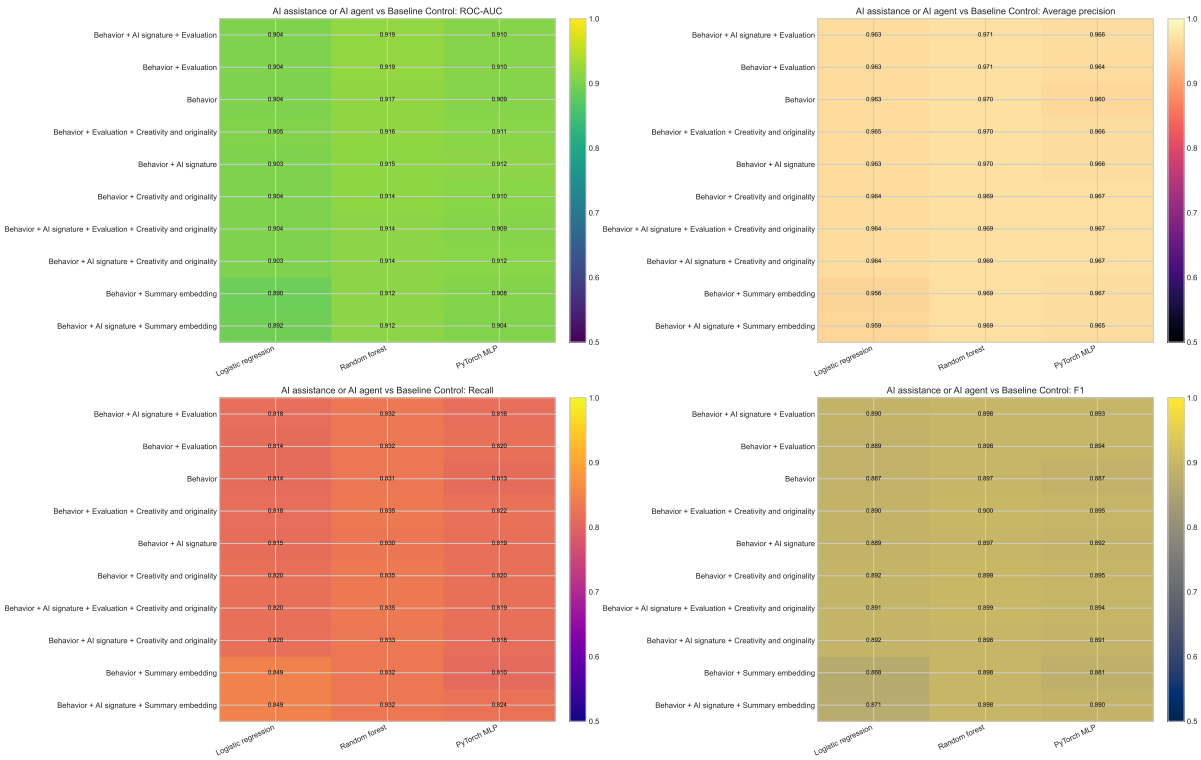

Figure 4: Top classifier combinations: AI assistance or AI agent vs. Baseline Control

Notes: The positive class combines human `AI assistance` responses and all AI agent responses; the comparison group is Baseline Control. The four panels report cross-validated ROC-AUC, average precision, recall, and F1. Rows show the top 10 feature-family combinations, ranked by classifier performance, and columns show the classifier used (logistic regression, random forest, and PyTorch MLP). Feature families are *Behavior*, *AI signature*, *Evaluation*, *Creativity and originality*, and *Summary embedding*. The summary embedding family is based on 384-dimensional summary-text embeddings from `sentence-transformers/all-MiniLM-L6-v2`. Larger values indicate better classifier performance. Sample sizes are $N = 412$ for Baseline Control and $N = 959$ for the combined AI group. The full heatmap with all feature-family combinations is reported in Appendix Figure A.9. See Appendix Figure A.7 for the classifier based only on the experimental results. Furthermore, for the broader comparison that also includes respondents randomized to low-quality AI assistance, see Appendix Figure A.8.

tive to peers. For survey researchers, the main takeaway is clear: effective authentication should not rely on a generic detector applied after fieldwork. It should combine complementary evidence about both process and product.

Researchers should build authentication workflows with two minimal characteristics. First, they should measure both how responses are produced and what the submitted text looks like. Process indicators such as paste attempts, keystrokes, deletions, pauses, and timing help reveal whether a response was drafted within the survey environment. Textual indicators such as AI-benchmark similarity, peer originality, and source alignment help reveal whether the final response is unusually close to machine-generated or otherwise standardized output. A workflow focused on only one of these dimensions will miss

important forms of AI use. Second, authentication workflows should be study-specific. Benchmark answers, thresholds, and classification rules should be tied to the exact prompt and task rather than imported from a general-purpose detector trained elsewhere. Our paper also suggests that synthetic AI agents can be useful during pretesting, helping researchers identify prompts that are especially easy to outsource before a survey is fielded.

Our findings also come with two important caveats. First, not all surveys will define problematic AI use in the same way. We therefore encourage researchers to decide ex ante what open-ended items are intended to measure and to align their authentication workflow with that substantive goal. Second, our approach does not imply a universal exclusion rule. It provides transparent evidence about likely AI involvement, but decisions about flagging, down-weighting, or exclusion remain design choices that should be justified in light of the research question and accompanied by sensitivity analyses.

Implementing the kind of workflow we identify here is a feasible way to improve the credibility of survey data. Yet researchers can do more. Future research could further explore the use of synthetic agents as a design-stage stress test. More broadly, the task for survey methodology is not simply to detect "AI text," but to design surveys that can document whether responses were produced by respondents themselves, under what conditions, and with what degree of human effort.

References

Argyle, Lisa P., Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. 2023. Out of One, Many: Using Language Models to Simulate Human Samples. *Political Analysis* 31 (3): 337–351.

Bell, Andrew M., and Thomas Gift. 2023. Fraud in Online Surveys: Evidence from a Nonprobability, Subpopulation Sample. *Journal of Experimental Political Science* 10 (1): 148–153.

Claassen, Joshua, Jan Karem Höhne, Ruben Bach, and Anna-Carolina Haensch. 2026. Identifying Bots Through LLM-Generated Text in Open Narrative Responses: A Proof-of-Concept Study. *Social Science Computer Review,* 08944393251408022.

Cohen, Denis. 2020. Between strategy and protest: how policy demand, political dissatisfaction and strategic incentives matter for far-right voting. *Political Science Research and Methods* 8 (4): 662–676.

Cohen, Mollie J., and Zach Warner. 2021. How to Get Better Survey Data More Efficiently. *Political Analysis* 29 (2): 121–138.

Crossley, Scott, Yu Tian, Joon Suh Choi, Langdon Holmes, and Wesley Morris. 2024. Plagiarism Detection Using Keystroke Logs, 476–483.

Dathathri, Sumanth, Abigail See, Sumedh Ghaisas, Po-Sen Huang, Rob McAdam, Johannes Welbl, Vandana Bachani, Alex Kaskasoli, Robert Stanforth, Tatiana Matejovicova, et al. 2024. Scalable watermarking for identifying large language model outputs. *Nature* 634 (8035): 818–823.

Deane, Paul, Mo Zhang, Jiangang Hao, and Chen Li. 2026. Using Keystroke Dynamics to Detect Nonoriginal Text. *Journal of Educational Measurement* 63 (1): e12431.

Diedenhofen, Birk, and Jochen Musch. 2017. PageFocus: Using paradata to detect and prevent cheating on online achievement tests. *Behavior Research Methods* 49 (4): 1444–1459.

Fuchs, Anna, Anna-Carolina Haensch, and Wiebke Weber. 2026. AI for Survey Design: Generating and Evaluating Survey Questions with Large Language Models. *SocArXiv.*

Heath, Anthony, Stephen Fisher, and Shawna Smith. 2005. THE GLOBALIZATION OF PUBLIC OPINION RESEARCH. *Annual Review of Political Science* 8 (1): 297–333.

Höhne, Jan Karem, Joshua Claassen, Saijal Shahania, and David Broneske. 2025. Bots in web survey interviews: A showcase. *International Journal of Market Research* 67 (1): 3–12.

Hullman, Jessica, David Broska, Huaman Sun, and Aaron Shaw. 2026. *This human study did not involve human subjects: Validating LLM simulations as behavioral evidence.*

Kane, John V., and Jason Barabas. 2019. No Harm in Checking: Using Factual Manipulation Checks to Assess Attentiveness in Experiments. *American Journal of Political Science* 63 (1): 234–249.

Kosmyna, Nataliya, Eugene Hauptmann, Ye Tong Yuan, Jessica Situ, Xian-Hao Liao, Ashly Vivian Beresnitzky, Iris Braunstein, and Pattie Maes. 2025. *Your Brain on ChatGPT: Accumulation of Cognitive Debt when Using an AI Assistant for Essay Writing Task.*

Lin, Gechun. 2025. Using cross-encoders to measure the similarity of short texts in political science. *American Journal of Political Science* 69 (4): 1600–1616.

Panizza, Folco, Yara Kyrychenko, and Jon Roozenbeek. 2026. Survey-taking AI tools surpass human abilities. Here's what we can do about it. *Nature* 650 (8101): 293–295.

Park, Joon Sung, Carolyn Q Zou, Aaron Shaw, Benjamin Mako Hill, Carrie J Cai, Meredith Ringel Morris, Robb Willer, Percy Liang, and Michael S Bernstein. 2025. Simulating Human Behavior with AI Agents.

Pinzón, Natalia, Vikram Koundinya, Ryan E. Galt, William O'R Dowling, Marcela Baukloh, Namah C. Taku-Forchu, Tracy Schohr, et al. 2024. AI-powered fraud and the erosion of online survey integrity: an analysis of 31 fraud detection strategies. *Frontiers in Research Metrics and Analytics* 9.

Rheault, Ludovic, and Christopher Cochrane. 2020. Word Embeddings for the Analysis of Ideological Placement in Parliamentary Corpora. *Political Analysis* 28 (1): 112–133.

Rodriguez, Pedro L., Arthur Spirling, and Brandon M. Stewart. 2023. Embedding Regression: Models for Context-Specific Description and Inference. *American Political Science Review* 117 (4): 1255–1274.

Sarstedt, Marko, Susanne J. Adler, Lea Rau, and Bernd Schmitt. 2024. Using large language models to generate silicon samples in consumer and marketing research: Challenges, opportunities, and guidelines. *Psychology & Marketing* 41 (6): 1254–1270.

Stagnaro, Michael Nicholas, James Druckman, Antonio Alonso Arechar, Robb Willer, DG Rand, et al. 2024. Representativeness versus attentiveness: a comparison across nine online survey samples.

Suguri Motoki, Fabio Yoshio, Januário Monteiro, Ricardo Malagueño, and Victor Rangel. 2024. *Large Language Models in Survey Research: Generating Synthetic Data and Unlocking New Possibilities.* SSRN Scholarly Paper. Rochester, NY.

Terry, John, Gerald Strait, Steve Alsarraf, Emily Weinmann, and Allison Waychoff. 2025. Artificial intelligence in scale development: evaluating AI-generated survey items against gold standard measures. *Current Psychology* 44 (20): 16339–16350.

Thomas, Kathrin. 2024. The Advent of Survey Experiments in Politics and International Relations. *Government and Opposition* 59 (1): 297–320.

Westwood, Sean J. n.d. The potential existential threat of large language models to online survey research. *Proceedings of the National Academy of Sciences of the United States of America* 122 (47): e2518075122.

Westwood, Sean J., and Samuel Frederick. n.d. Reply to Van der Stigchel et al.: Empirical evidence that AI survey contamination is real and substantial. *Proceedings of the National Academy of Sciences of the United States of America* 123 (8): e2537420123.

Yan, Duanli, Michael Fauss, Jiangang Hao, and Wenju Cui. n.d. Detection of AI-generated Essays in Writing Assessments.

Zhang, Simone, Janet Xu, and Aj Alvero. 2025. Generative AI Meets Open-Ended Survey Responses: Research Participant Use of AI and Homogenization. *Sociological Methods & Research* 54 (3): 1197–1242.

Ziegler, Jeffrey. 2022. A Text-As-Data Approach for Using Open-Ended Responses as Manipulation Checks. *Political Analysis* 30 (2): 289–297.

A Online Appendix

Contents

A.1 Extended Related Work 1
A.2 Experimental Design 2
A.3 An Integrated Toolkit for AI-Response Detection 6
A.3.1 Front-End Manipulation: Copy-Paste Disablement 6
A.3.2 Behavioral Monitoring: Keystroke Dynamics Logging 7
A.3.3 Back-End Linguistic Analysis: AI Similarity and Originality Checks 8
A.4 Data Collected 9
A.5 A Synthetic AI Agent for Pre-Launch Survey Vulnerability Audits 10
A.6 Tables 13
A.7 Additional Results 13
A.8 Additional Metrics Classification 13

A.1 Extended Related Work

Recent discussions in survey methodology and political methodology converge on a common point: credible online data collection now depends less on any single screening device than on transparent workflows that document how suspicious cases are identified and handled. Reviews of survey fraud show that the usable yield of low-barrier web recruitment has deteriorated sharply and that standard defensive batteries are readily bypassed (Pinzón et al. 2024). Work in political methodology similarly shows that explicit combinations of automated and human-coded checks can recover much of the data quality achieved by intensive supervision (Cohen and Warner 2021), while studies of targeted populations document how adversarial respondents learn to defeat simple screeners (Bell and Gift 2023). The lesson is not simply that online data are noisy. It is that valid inference increasingly depends on observable, reviewable quality-control procedures.

AI assistance sharpens that requirement because it introduces a different kind of measurement failure. One useful distinction is between deception and delegation: sometimes no real respondent stands behind the record, and sometimes a real respondent is present but outsources part of the task to AI. These cases call for overlapping, but not identical, responses. Classic fraud and satisficing often degrade answers in visible ways, such as through speeding, straightlining, contradiction, or nonsense. AI-assisted delegation can do the opposite: it can improve the visible product while weakening the link between the response and the respondent's own reasoning. That matters for the open-ended tasks political scientists use to evaluate learning, comprehension, attribution, or justification. The measurement problem is therefore not only whether a respondent is human, but also whether the observed answer was produced under conditions consistent with the construct being measured (Zhang, Xu, and Alvero 2025; Westwood, n.d.; Westwood and Frederick, n.d.). Not every instance of AI use is equally damaging, but when a task is meant to capture unaided cognition, authorship becomes part of measurement.

The emerging literature has made important progress, but most contributions emphasize one part of the problem:

- **Survey fraud and data-quality audits:** This work shows that low-barrier online samples are increasingly vulnerable and that simple defensive batteries are often ineffective. Its main contribution is to document the scale of the threat and the need for layered quality control (Pinzón et al. 2024; Bell and Gift 2023).

- **Behavioral and paradata-based detection:** These studies focus on process indicators such as response speed, keystroke dynamics, paste behavior, and other paradata to identify copied or nonoriginal text. They show that writing traces can be highly informative, especially when the response was transferred or transcribed rather than composed in the survey box (Westwood and Frederick, n.d.; Deane et al. 2026; Diedenhofen and Musch 2017; Crossley et al. 2024).

- **Text-based or classifier-based AI detection:** This line of work tries to distinguish human from machine-generated responses using linguistic features. These tools can perform well in bounded settings, but often depend on the prompt, corpus, or model family used for training, which raises concerns about portability across survey contexts (Claassen et al. 2026; Zhang, Xu, and Alvero 2025).

- **Agent-vulnerability studies:** Recent work shows that autonomous AI respondents

can pass standard screeners, maintain coherent personas, and generate plausible answers at scale. The key contribution here is to demonstrate that the threat is no longer hypothetical and that standard attention checks are not designed for this kind of contamination (Westwood, n.d.; Panizza, Kyrychenko, and Roozenbeek 2026).

- **LLMs as substitutes for respondents or survey-design tools:** An adjacent literature asks when LLMs can stand in for human respondents, generate synthetic samples, or help evaluate survey items. Its focus is calibration, substitution, and design assistance rather than authentication of human-authored responses (Hullman et al. 2026; Argyle et al. 2023; Suguri Motoki et al. 2024; Fuchs, Haensch, and Weber 2026; Terry et al. 2025).

Our contribution differs from these approaches in several ways. First, we do not study only whether a response *looks* machine generated; we study whether it was produced under authorship conditions consistent with the construct the survey is meant to capture. Second, we do not rely on a single family of signals. Instead, we combine preventive design, process paradata, prompt-specific AI benchmarks, peer-distinctiveness measures, evaluator-based assessments, and, when applicable, source-text alignment in one auditable framework. Third, our approach is explicitly task-specific rather than a universal black-box detector: similarity is evaluated against benchmark AI output generated for the exact prompt, article, and instructions used in the study. Fourth, we validate the framework using experimentally assigned and directly observed AI access embedded inside the survey environment, rather than relying only on post hoc suspicion. Finally, we use a synthetic AI agent not as a substitute for human subjects, but as a benchmark and prefield auditing tool that helps diagnose survey vulnerability before launch.

For these reasons, the paper's contribution is broader than a better detector of AI text. It provides a survey-embedded authentication workflow for researchers who need to document not only the final product of responding, but also the conditions under which that product was generated.

A.2 Experimental Design

This study validates a survey-embedded approach to detecting AI-assisted responding in open-ended tasks. The design is structured to resemble a realistic paid task environment while allowing us to observe both the final responses and the process through which they are produced.

Recruitment and cover story. Participants were recruited from an online labor platform (CloudResearch Connect) and told they were completing a short paid writing task as part of a project to improve written materials. The task was framed as routine editing work rather than as a study about AI. This framing reduces demand effects and discourages respondents from strategically altering their behavior based on the study's purpose.

Task environment. All participants completed two tasks commonly found in online work: (i) summarizing a short article (approximately 150 words) and (ii) proposing a concise title. Instructions emphasized clarity, accuracy, and the use of participants' own

words. After submission, participants answered a comprehension question to ensure engagement with the source text.

To avoid confounding by content, participants were randomly assigned one of several short factual passages covering neutral topics. This ensures that differences across conditions are not driven by a specific article.

Treatment conditions and framing. The experimental manipulation varies whether participants have access to an AI writing assistant and how that access is described.

In the control condition, participants were explicitly told that external AI tools were not allowed and that submissions would be screened for AI use. This restriction was reinforced by disabling copy-paste functionality within the survey interface.

In the AI-access conditions, participants were told that a writing support tool was available to help them complete the task. The tool was presented as part of an effort to improve workflow efficiency and assist workers in producing higher-quality output. This framing mirrors how such tools are commonly introduced in real work environments.

In one condition, the tool was additionally described as part of a longer-term effort to automate similar tasks. This variation allows us to distinguish between immediate exposure to AI assistance and awareness of its broader implications.

Embedded AI tool. Rather than allowing respondents to rely on external AI platforms, we built a custom AI assistant for this study and embedded it directly into the task environment. The tool was implemented as a Streamlit application and connected to OpenAI's `gpt-4o-mini` model through the API. Embedding the assistant in this way allowed respondents to access AI support without leaving the survey workflow, while also allowing us to observe and log their interactions with the tool.

Participants in the AI conditions could paste the article into the application and receive an initial summary. They could then request up to three rounds of revisions by typing additional instructions, with the model returning an updated version after each round. This setup was designed to resemble a realistic writing-assistance tool while keeping all AI use inside a controlled interface.

The visible interface was kept constant across AI conditions. What changed across conditions was the prompt used to govern the model's behavior. In the regular-quality condition, `gpt-4o-mini` was instructed to generate a summary of about 150 words that reflected the core elements emphasized in the task instructions, including the main idea of the text, the key information, and its significance. The model was also instructed to remain faithful to the source text and to rely on its own wording rather than reproducing the article verbatim. In the low-quality condition, the same model was prompted to produce deliberately weak assistance: shorter and less complete summaries, inconsistent style, omission of central points, and occasional irrelevant or inaccurate details. Participants in both conditions could still request revisions, but in the low-quality condition the model continued to respond according to the degraded prompt. Figures A.1a and A.1b present examples of the two versions.

Embedding the tool directly in the study served an important methodological purpose. It created a treatment in which AI assistance was available as part of the task environment rather than through an external and unobserved platform, and it allowed us to collect detailed logs of tool use. At the session level, we recorded the article text pasted into

the app, the model's outputs, respondents' revision requests, and basic timing metadata. This makes it possible to distinguish assignment to AI access from actual interaction with the tool and to analyze not only the final response submitted by the participant, but also the process through which AI assistance was used.

Thus, a key feature of the design is that AI assistance is *observed* rather than inferred. Because the tool is embedded within the survey, we record all interactions, including the text entered, the model's outputs, and revision requests. This allows us to distinguish between access to AI and actual usage behavior.

Behavioral instrumentation. The survey interface was instrumented to capture detailed process data, including keystrokes, deletions, pauses, paste attempts, and timing. These measures provide direct evidence on how responses are produced. Human-authored responses typically involve iterative drafting and revision, whereas AI-assisted responses tend to exhibit more linear production patterns or signs of transfer.

Outcome logic. The design allows us to compare both the content of responses and the process that generated them. On the textual side, we evaluate similarity to prompt-specific AI benchmarks, alignment with the source article, and distinctiveness relative to peer responses. On the behavioral side, we examine indicators such as keystroke activity, editing intensity, and production time.

Together, these measures allow us to identify whether AI assistance leaves a consistent and detectable signature, even when the final response appears fluent and relevant.

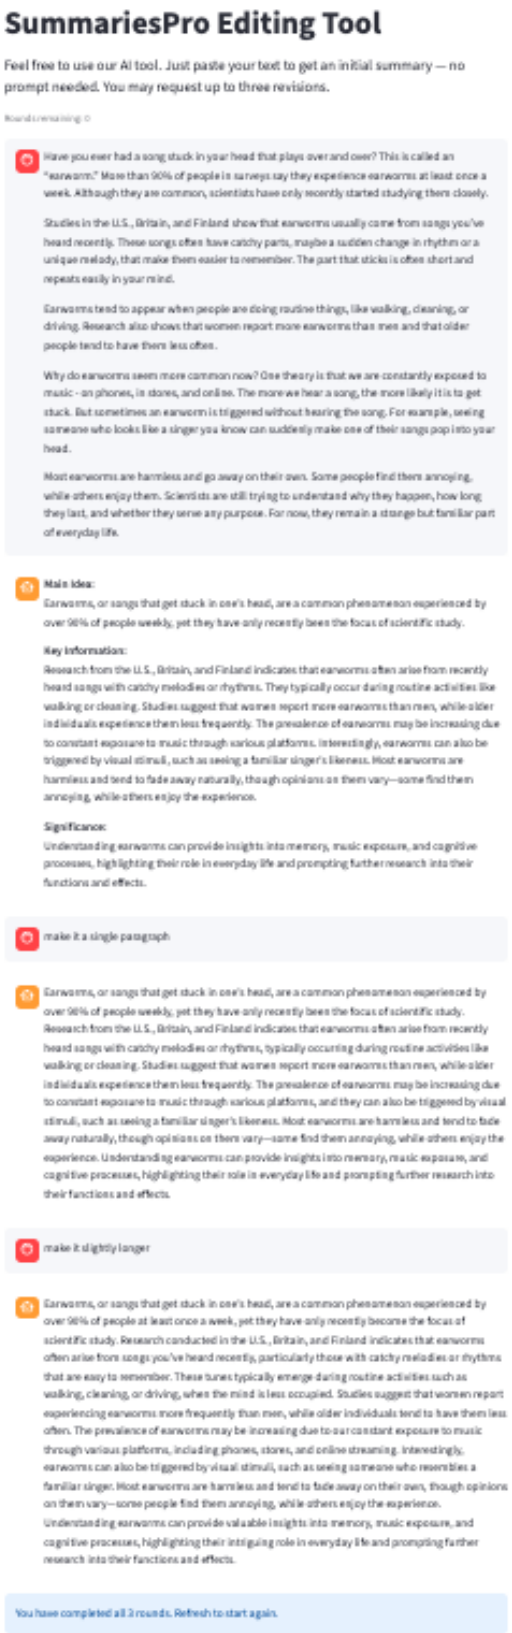

(a) Regular Quality

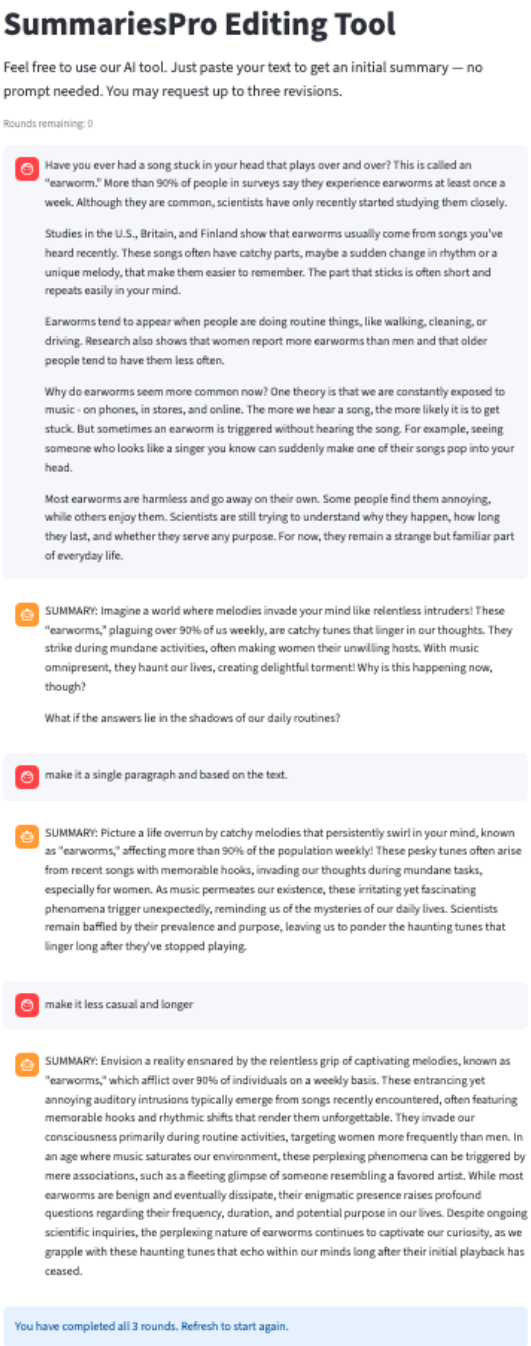

(b) Low Quality

Figure A.1: Examples of interaction with the Embedded AI App across conditions

A.3 An Integrated Toolkit for AI-Response Detection

Our proposed toolkit addresses AI-assisted responding on three levels: (1) front-end deterrence, (2) real-time behavioral logging, and (3) post-hoc linguistic analysis. By combining these, we create a robust composite measure of "AI involvement" in a given open-ended response. Here we detail each component and the rationale behind it.

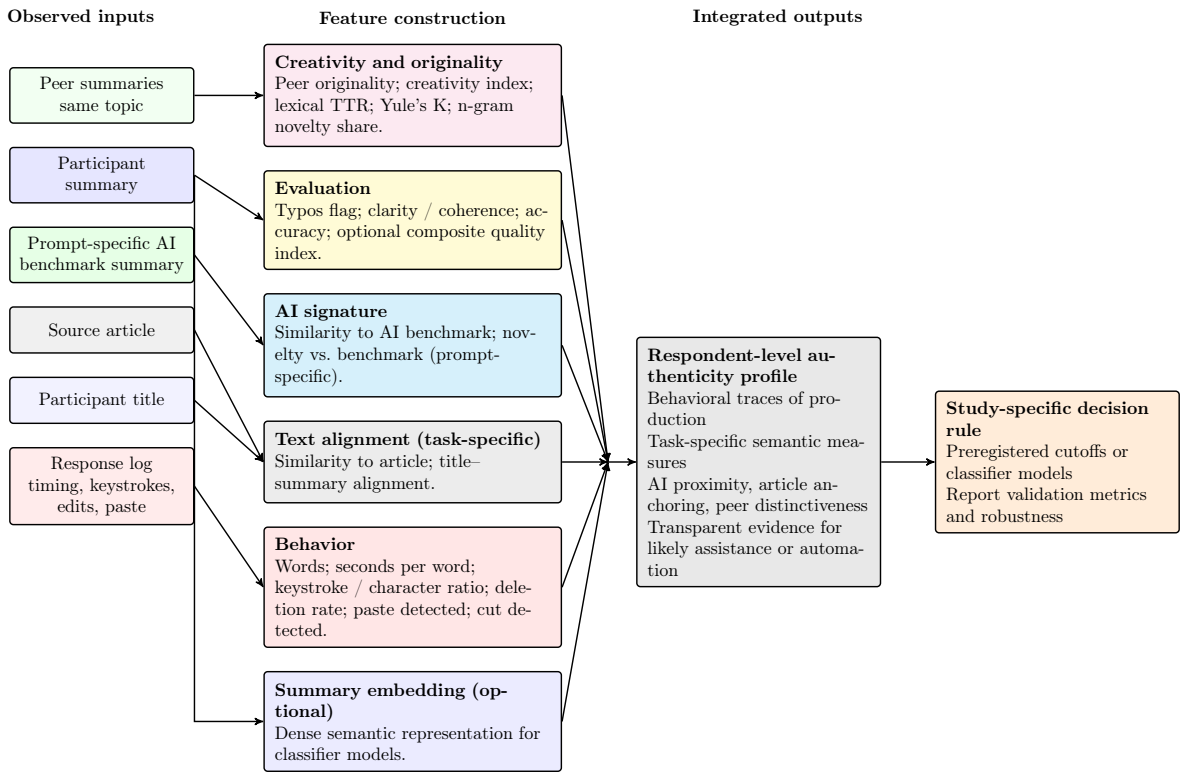

Figure A.2: Workflow for the open-ended authentication toolkit. The pipeline maps observed survey inputs into feature families—Behavior, Text alignment, AI signature, Evaluation, Creativity and originality, and an optional embedding branch—and combines them into a respondent-level authenticity profile used for study-specific classification or threshold-based flagging.

A.3.1 Front-End Manipulation: Copy-Paste Disablement

The first tool is a simple front-end intervention: disabling copy-paste functionality for certain survey questions. The logic is straightforward. A respondent leaning on an AI like ChatGPT typically would copy the AI's answer from another window and paste it into the survey text box. By technically blocking that action (via JavaScript embedded in the survey), we force respondents to manually type their answers in real time. This creates friction for those attempting to use external text generators. If a participant still chooses to use AI, they must transcribe the AI-generated text keystroke by keystroke, a slower and more laborious process that our system can then detect through behavioral cues.

This manipulation serves two purposes: deterrence and detection. On one hand, it likely dissuades some users from cheating; if copy-paste is unavailable, the convenience

of using AI drops, and more participants may answer in their own words. On the other hand, for those who persist in using AI (by re-typing an AI-provided answer), the act of doing so produces distinctive patterns we can capture.

In our implementation, if a paste is attempted, the script silently blocks it (nothing is inserted) and logs the event. The survey interface and question wording remain the same, preserving the user experience except for the disabled paste option. By being covert, this approach avoids alerting respondents to the specific detection mechanism (which could invite attempts to game it), while still upholding ethical standards – we do not prevent use of AI outright. Still, we ensure that if it is used, it becomes behaviorally evident.

A.3.2 Behavioral Monitoring: Keystroke Dynamics Logging

A cornerstone of our toolkit is the monitoring of typing behavior, or keystroke dynamics, during the answering process. We instrumented the survey text boxes with a logging script that records each key press, including timing, key codes, and whether it was an input or a deletion. These detailed paradata capture the process of composition for every open-ended response. From these logs, we derive features that can reveal whether an answer was composed through careful thought or generated elsewhere and transcribed.

Our approach builds on the idea that human writing is iterative, while AI-generated text (when transcribed by a human or pasted) appears in a more linear burst. Cognitive writing theory holds that humans plan, type, pause, revise, and edit in cycles; this manifests as starts and stops, deletions, and rephrasings in the keystroke record (Crossley et al. 2024). In contrast, if someone is copying an AI answer, the keystrokes tend to flow continuously as they transcribe or, in the case of a direct paste, a large chunk of text appears near-instantly with minimal backtracking.

Concretely, we log metrics such as: total keystrokes pressed, characters typed, backspace (delete) usage, pauses between keystrokes, and paste events. From these, we compute summary indicators. For example, a key metric is the ratio of characters to keystrokes (or its inverse). A genuinely composed answer typically involves many more keystrokes than the final character count (due to corrections and edits), yielding a low characters-per-keystroke ratio. An AI-generated answer that is pasted (or typed with almost no editing) will have a much higher ratio – possibly approaching 1.0 if the answer was typed out flawlessly or infinitely large if pasted in one go. We also look at "extra" keystrokes, defined as the total keystrokes minus the final number of characters in the submitted answer. This captures how much editing was done. A large number of extra keystrokes (relative to answer length) indicates the user spent effort revising the text (a human-like process), whereas near-zero extra keystrokes suggests the text arrived fully formed (suspicious of AI transcription).

Prior research strongly supports the efficacy of these indicators. Crossley et al. 2024 use keystroke logging to distinguish student essays written authentically versus copied from an outside source (e.g., an AI). They find that authentic writing showed frequent pauses, insertions, and deletions, whereas transcribed text appeared in longer uninterrupted bursts. Their detection model achieved 99% accuracy in classifying an essay as original vs. copied based on features like pause frequency and revision counts. We leverage the same insight: how the text was generated is important information. Our toolkit flags responses with anomalous writing dynamics.

A.3.3 Back-End Linguistic Analysis: AI Similarity and Originality Checks

The third component of the toolkit analyzes the textual content of responses after submission, using modern NLP methods to evaluate their similarity to AI-generated text and their degree of originality. While behavioral data capture how a response was written, the linguistic layer provides independent evidence about what was written. By integrating both, we obtain a more reliable signal of whether an answer reflects genuine human reasoning or AI assistance.

LLM Similarity For each survey question, we generate reference answers using a large language model (GPT-3.5 or GPT-4) under the same prompt presented to participants. Each participant's text and the corresponding AI reference are embedded into a high-dimensional semantic space using the text-embedding-3-small model. The cosine similarity between these embeddings provides a continuous measure of semantic proximity to the AI baseline. Higher similarity indicates that the response resembles what an LLM would produce, whereas lower similarity reflects a more distinctive and idiosyncratic phrasing. We interpret this score as an indicator of "LLM-likeness." Although similarity alone cannot prove AI use, unusually high correspondence between a participant's summary and the model's output suggests reliance on or imitation of AI-generated content.

Originality and Faithfulness To complement this measure of LLM similarity, we compute several indicators of originality and topic adherence based on embeddings and lexical properties. Each response is compared to the article it summarizes, yielding a faithfulness score that reflects how closely the participant's writing captures the source material's main ideas. Simultaneously, the novelty of the response is measured as its cosine distance from the AI baseline embedding, indicating the extent to which the participant's phrasing diverges from model-generated language. A related alignment metric captures the coherence between each participant's title and summary by measuring their cosine similarity within the same embedding space.

We further examine textual originality through surface-level metrics. Lexical richness is assessed using the type–token ratio and Yule's K statistic, which together capture the diversity and repetition of vocabulary. We also compute the share of bigrams and trigrams in a response that do not appear in the corresponding article—an indicator of unique phrasing and paraphrasing effort (Kosmyna et al. 2025). To assess variation across participants exposed to the same article, we calculate peer originality, defined as the embedding distance between a participant's summary and the centroid of all summaries for that topic. Responses that are semantically distinct from the group mean, while remaining faithful to the article, are interpreted as more original.

Each of these features is standardized within topic to control for differences in article complexity and then combined into a composite Creativity Index. The index weights novelty most heavily but also rewards faithfulness, internal coherence, and lexical variety. High values thus characterize responses that are both on-topic and written in the participant's own words, whereas low values indicate close resemblance to AI-generated or source text.

AI-as-Judge Evaluation In addition to quantitative similarity measures, we employ large language models as evaluators—or "AI judges"—to provide holistic assessments of the written responses. Each summary is presented to an LLM alongside the reference article, and the model produces ratings for clarity and coherence, factual accuracy, typo detection, and the degree to which the response appears to be written in the respondent's own words. This procedure mirrors recent advances in automated grading and text evaluation, where LLMs are used as consistent, low-variance evaluators across large datasets (Kosmyna et al. 2025). Although these judgments are not used as definitive classifiers of AI use, they serve as complementary indicators that capture dimensions of writing quality and originality that are difficult to express through embeddings or lexical statistics alone.

By combining these linguistic indicators with the behavioral data described earlier, the toolkit provides a multidimensional assessment of response authenticity. Rather than issuing a binary classification, it generates a detailed profile—capturing AI similarity, originality, and textual diversity—that allows researchers to identify suspicious cases and calibrate detection thresholds while minimizing false positives.

A.4 Data Collected

For each participant's summary response, we recorded the full text of the answer and a detailed keystroke log (including timestamps for each key event). Using the toolkit's back-end analysis, we generated a set of behavioral and linguistic indicators for every response. These measures jointly capture the process of composition (through keystroke dynamics) and the content of the written text (through embedding-based and lexical analysis). Specifically, we computed:

- **Character Count:** Total number of characters in the final submitted response.
- **Keystroke Count:** Total number of keys pressed, including letters, numbers, spaces, deletions, and other editing keys.
- **Extra Keystrokes:** Difference between keystroke count and character count, indicating the amount of revision and editing (extra keystrokes = keystrokes − characters).
- **Keystroke–Character Ratio:** Ratio of total key presses to final characters (keystrokes/characters), representing the extent of typing activity per retained character.
- **Paste Event:** Binary indicator denoting whether a paste action (`Ctrl/Command+V`) occurred during text entry.
- **Cosine Similarity to AI Answer:** Semantic similarity between the participant's response and a GPT-3.5–generated reference summary, computed using embeddings from the `text-embedding-3-small` model.
- **Cosine Similarity to Article:** Similarity between the participant's response and the original article text, reflecting topical faithfulness.

- **Novelty vs. AI:** Cosine distance ($1 - \text{cosine similarity}$) from the AI baseline, measuring how distinct the participant's phrasing and ideas are from model-generated language.
- **Summary–Title Alignment:** Cosine similarity between a participant's summary and title embeddings, used as a measure of internal coherence.
- **Lexical Diversity:** Vocabulary richness estimated via the type–token ratio (TTR) and Yule's K statistic.
- **N-gram Novelty:** Proportion of bigrams and trigrams in the participant's response not present in the original article, capturing unique phrasing.
- **Peer Originality:** Embedding distance between a participant's summary and the centroid of all summaries written about the same article, indicating distinctiveness among peers.
- **AI-as-Judge Evaluation:** Structured ratings provided by a large language model (GPT-4) assessing clarity, coherence, accuracy, and typos.

All variables are standardized within topic to account for differences in article complexity, and they are combined in subsequent analyses to construct composite indicators such as the *Creativity Index* and overall measures of AI-likeness and response authenticity.

These measures were collected uniformly for all responses. The keystroke features and content features were then analyzed to see which best separate the two experimental groups (which serve as a proxy for likely AI use, since presumably AI-Allowed should have a higher incidence of AI-written answers). Because we experimentally manipulated one factor, we can interpret differences more cleanly than in an observational setting. In practice, a researcher using this toolkit in a real-world survey (without experimental conditions) would rely on a combination of metrics and perhaps a trained model to flag cases of AI-generated content.

A.5 A Synthetic AI Agent for Pre-Launch Survey Vulnerability Audits

An AI agent, in this setting, is a bounded software system that can perceive a survey page, represent the current task, generate answers conditional on a stable persona, and execute those answers in a browser without manual intervention. Rather than treating AI exposure as an abstract risk, researchers can use such an agent as a pre-launch stress test on their own instrument. If a controlled synthetic respondent can complete the survey reliably, answer comprehension items from on-screen context, and produce plausible open-ended text at low cost, then the instrument is vulnerable to AI-mediated responding. The implementation described here was built for controlled benchmarking on an owned survey. Although the prototype targeted standard Qualtrics page structures, the logic is portable to other web surveys that expose ordinary HTML form elements.

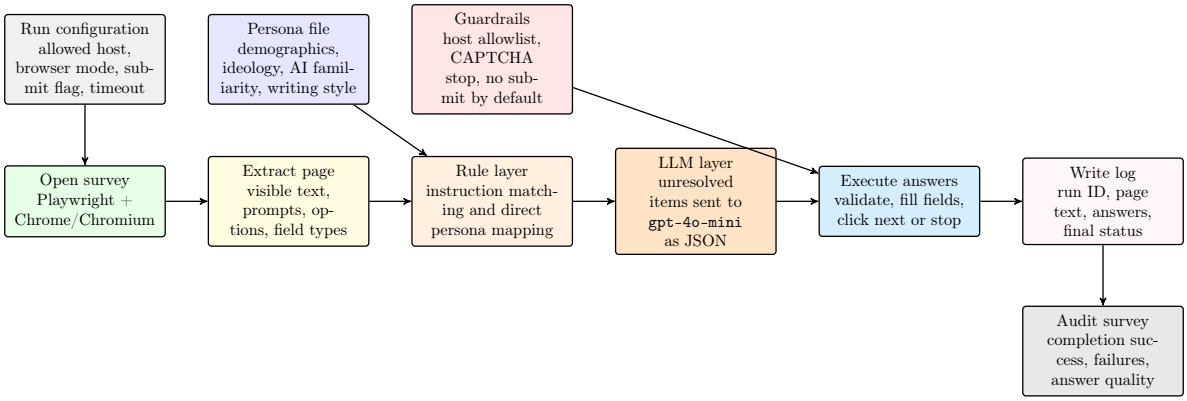

Figure A.3: Synthetic-agent audit workflow. The agent opens the survey in a controlled browser, extracts visible questions into structured fields, answers simple items with rules, sends unresolved items to `gpt-4o-mini`, executes the validated answers in the browser, and stores a full run log for pre-launch auditing.

Using the agent as a template for vulnerability detection. Researchers can use this design as a pre-launch audit protocol. First, create a small pool of realistic personas rather than a single synthetic respondent. Second, run the instrument in dry mode and inspect the logs for unsupported widgets, broken parsers, or pages that the agent cannot interpret. Third, run a small submitted batch on the pilot version of the owned survey. Fourth, inspect whether the agent can pass attention checks, answer comprehension items from on-screen text, and generate plausible open-ended responses. Finally, compare completion success, answer diversity across personas, and the substantive quality of the generated text. A survey that a low-cost agent can complete reliably under these conditions should be treated as vulnerable to AI-mediated responding before field launch, even if ordinary fraud checks would not flag those cases.

Model and runtime. The benchmark agent was implemented in Python and used the OpenAI API together with Playwright browser automation. In the configuration used for the audit, the answer model was `gpt-4o-mini`, temperature was set to 0.5, the browser ran headless by default, the timeout was 15,000 milliseconds, and the agent retained up to four prior pages of textual context when answering later items. The core software stack was: `openai>=1.0.0` for model calls, `playwright==1.58.2` for browser control, Chrome/Chromium for page execution, JSON persona and configuration files for controlled variation, and JSON log files for post-run auditing. For larger batches, the same runner could be containerized and sharded across cloud jobs, with optional artifact upload through Google Cloud Storage.

What the agent actually received. The model did not receive unconstrained browsing instructions. Instead, it received a structured payload containing: (i) persona fields such as age group, gender, education, ethnicity, work status, occupation, ideology, AI familiarity, AI usage, description, and writing style; (ii) recent page text from earlier screens; (iii) current page text; (iv) prior answers; and (v) a list of unresolved fields with the prompt, field type, available options, and local question context. This design matters

because it makes the agent auditable: researchers can see exactly what information the model had at each step.

Core instruction prompt. The system prompt supplied to the model was intentionally narrow and operational rather than conversational:

> "You are generating synthetic survey responses. Answer consistently as the provided persona. Follow the question instructions exactly. Return only valid JSON."

The response schema was similarly constrained:

```
{"answers": {"FIELD_ID": "string or list of strings"}}
```

Hybrid answer strategy. Not every answer was generated freely by the language model. Before calling the model, the agent first applied deterministic rules. If a question explicitly instructed the respondent to select a quoted option, the agent matched that option directly. If a field asked for demographics or background items represented in the persona file, the agent mapped the stored persona values onto the available response options. Only unresolved items were passed to the model. After the model responded, every returned answer was validated against the extracted response options before being entered into the survey. This hybrid design reduced unnecessary model calls and made the agent more consistent across repeated runs.

Survey parsing and supported question types. The page parser inspected visible question blocks and transformed them into a structured representation containing a field identifier, question identifier, prompt text, response type, required status, selector, response options, and local context. In the implementation used here, the agent supported five common classes of items: single-choice questions, multi-select questions, dropdown menus, text-entry boxes, and radio-matrix rows. Unsupported widgets were not silently guessed; instead, they appeared as failures or incomplete pages in the run log, which is itself useful information for a pre-launch vulnerability audit.

How the agent was run. A dry run opened the survey, answered pages, and stopped before the final submission button:

```
python -m survey_agent \
  --survey-url "YOUR_SURVEY_URL" \
  --persona-file configs/persona.json \
  --config-file configs/agent_config.json
```

Adding `-submit` allowed final submission on an owned survey. Adding `-count`, `-persona-pool-file`, and `-sleep-seconds` enabled small benchmark batches across multiple synthetic personas. A visible browser mode was also available through `-headful`, and a manual pause could be triggered at a CAPTCHA gate if the researcher wanted to inspect that failure mode.

Guardrails and audit value. Each run wrote a JSON log containing the final status, extracted page text, parsed fields, chosen answers, and error information when relevant. These safeguards are methodologically useful because the goal is not covert completion of arbitrary surveys, but controlled benchmarking of one's own instrument before deployment.

A.6 Tables

Table A.1: Summary-only coefficient regressions: AI assistance (AI_only) vs Control (refers to Figure 2)

	Dependent variable												
	Key/char	Delete rate	Paste	Words	Sec./word	Cut	AI sim.	Typos	Peer orig.	Creativ. index	Yule's K	TTR	N-gram nov.
AI assistance	-1.040***	-0.790***	0.813***	0.396***	-0.186***	-0.056	0.887***	-0.533***	-0.706***	-0.827***	-0.277***	0.148**	-0.331***
	(0.060)	(0.064)	(0.064)	(0.068)	(0.071)	(0.071)	(0.059)	(0.067)	(0.065)	(0.063)	(0.067)	(0.067)	(0.067)
Topic: Dreams	-0.055	-0.093	0.072	0.234*	-0.116	0.114	0.526***	-0.023	0.093	-0.011	0.016	-0.091	-0.316**
	(0.107)	(0.112)	(0.122)	(0.125)	(0.169)	(0.112)	(0.110)	(0.128)	(0.120)	(0.111)	(0.112)	(0.118)	(0.153)
Topic: Kindness	-0.000	-0.020	0.006	0.147	-0.082	-0.000	0.272**	0.068	0.110	0.012	0.472***	-0.468***	0.291**
	(0.106)	(0.117)	(0.115)	(0.139)	(0.164)	(0.004)	(0.107)	(0.125)	(0.123)	(0.115)	(0.144)	(0.134)	(0.136)
Topic: Music	0.142	0.153	-0.049	-0.108	-0.039	0.007	-0.157	0.248*	0.135	0.192	-0.308**	0.338**	0.108
	(0.113)	(0.125)	(0.118)	(0.125)	(0.189)	(0.009)	(0.117)	(0.131)	(0.130)	(0.122)	(0.121)	(0.131)	(0.132)
Topic: Sleep	0.036	-0.012	0.119	0.060	0.016	0.004	0.446***	0.057	0.096	-0.062	-0.298***	0.257**	0.013
	(0.105)	(0.116)	(0.117)	(0.116)	(0.178)	(0.006)	(0.126)	(0.127)	(0.120)	(0.123)	(0.090)	(0.123)	(0.147)
Topic: Smelling	0.103	0.085	-0.024	-0.026	-0.225	0.109	-0.042	-0.096	0.192	0.131	-0.030	0.109	0.508***
	(0.107)	(0.121)	(0.113)	(0.122)	(0.164)	(0.108)	(0.090)	(0.128)	(0.118)	(0.106)	(0.105)	(0.126)	(0.123)
Topic: Sugar	0.119	0.118	0.077	0.078	0.068	0.364**	-0.080	0.059	0.138	0.143	0.305***	-0.198	0.337***
	(0.116)	(0.126)	(0.114)	(0.126)	(0.171)	(0.179)	(0.090)	(0.123)	(0.120)	(0.109)	(0.111)	(0.129)	(0.124)
Constant	0.472***	0.363***	-0.439***	-0.255***	0.143	-0.060*	-0.583***	0.221**	0.246***	0.358***	0.108	-0.059	0.027
	(0.077)	(0.093)	(0.076)	(0.089)	(0.137)	(0.032)	(0.069)	(0.097)	(0.085)	(0.077)	(0.078)	(0.096)	(0.100)
N	828	828	828	828	828	828	828	828	828	828	828	828	828
Positive class N	416	416	416	416	416	416	416	416	416	416	416	416	416
Control N	412	412	412	412	412	412	412	412	412	412	412	412	412
R^2	0.275	0.162	0.169	0.050	0.017	0.017	0.265	0.078	0.127	0.178	0.098	0.075	0.092
Adj. R^2	0.268	0.155	0.162	0.042	0.008	0.009	0.258	0.070	0.120	0.171	0.090	0.067	0.085
Topic FE	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes

Notes: Standard errors in parentheses; HC1 robust SEs. All regressions include topic fixed effects, standardized outcomes, and omit the Coloring topic. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

A.7 Additional Results

Figure A.4 reports the extended benchmark coefficient plot for the original human sample, comparing each AI treatment arm separately to the same baseline control group. Positive coefficients indicate that the corresponding outcome is higher in the given AI arm than in the control group, while negative coefficients indicate lower values relative to control.

A.8 Additional Metrics Classification

Table A.2: Summary-only coefficient regressions: AI assistance and AI agent vs Baseline Control (refers to Figure 3)

Panel A. AI assistance vs Baseline Control

	Dependent variable												
	Sec./word	Words	Delete rate	Key/char	Paste	Cut	AI sim.	Typos	TTR	Yule's K	N-gram nov.	Peer orig.	Creativ. index
AI assistance	-0.186***	0.396***	-0.790***	-1.040***	0.813***	-0.056	0.887***	-0.533***	0.148**	-0.277***	-0.331***	-0.810***	-0.892***
	(0.071)	(0.068)	(0.064)	(0.060)	(0.064)	(0.071)	(0.059)	(0.067)	(0.067)	(0.067)	(0.067)	(0.064)	(0.061)
Topic: Dreams	-0.116	0.234*	-0.093	-0.055	0.072	0.114	0.526***	-0.023	-0.091	0.016	-0.316**	0.012	-0.073
	(0.169)	(0.125)	(0.112)	(0.107)	(0.122)	(0.112)	(0.110)	(0.128)	(0.118)	(0.112)	(0.153)	(0.118)	(0.113)
Topic: Kindness	-0.082	0.147	-0.020	-0.000	0.006	-0.000	0.272**	0.068	-0.468***	0.472***	0.291**	0.046	0.027
	(0.164)	(0.139)	(0.117)	(0.106)	(0.115)	(0.004)	(0.107)	(0.125)	(0.134)	(0.144)	(0.136)	(0.120)	(0.115)
Topic: Music	-0.039	-0.108	0.153	0.142	-0.049	0.007	-0.157	0.248*	0.338**	-0.308**	0.108	0.091	0.038
	(0.189)	(0.125)	(0.125)	(0.113)	(0.118)	(0.009)	(0.117)	(0.131)	(0.131)	(0.121)	(0.132)	(0.127)	(0.129)
Topic: Sleep	0.016	0.060	-0.012	0.036	0.119	0.004	0.446***	0.057	0.257**	-0.298***	0.013	0.093	-0.118
	(0.178)	(0.116)	(0.116)	(0.105)	(0.117)	(0.006)	(0.126)	(0.127)	(0.123)	(0.090)	(0.147)	(0.116)	(0.119)
Topic: Smelling	-0.225	-0.026	0.085	0.103	-0.024	0.109	-0.042	-0.096	0.109	-0.030	0.508***	0.119	0.078
	(0.164)	(0.122)	(0.121)	(0.107)	(0.113)	(0.108)	(0.090)	(0.128)	(0.126)	(0.105)	(0.123)	(0.119)	(0.104)
Topic: Sugar	0.068	0.078	0.118	0.119	0.077	0.364**	-0.080	0.059	-0.198	0.305***	0.337***	0.111	0.169
	(0.171)	(0.126)	(0.126)	(0.116)	(0.114)	(0.179)	(0.090)	(0.123)	(0.129)	(0.111)	(0.124)	(0.116)	(0.105)
Constant	0.143	-0.255***	0.363***	0.472***	-0.439***	-0.060*	-0.583***	0.221**	-0.059	0.108	0.027	0.339***	0.429***
	(0.137)	(0.089)	(0.093)	(0.077)	(0.076)	(0.032)	(0.069)	(0.097)	(0.096)	(0.078)	(0.100)	(0.082)	(0.077)
N	828	828	828	828	828	828	828	828	828	828	828	828	828
Positive class N	416	416	416	416	416	416	416	416	416	416	416	416	416
Control N	412	412	412	412	412	412	412	412	412	412	412	412	412
R^2	0.017	0.050	0.162	0.275	0.169	0.017	0.265	0.078	0.075	0.098	0.092	0.166	0.210
Adj. R^2	0.008	0.042	0.155	0.268	0.162	0.009	0.258	0.070	0.067	0.090	0.085	0.159	0.203
Topic FE	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes

Panel B. AI agent vs Baseline Control

	Dependent variable												
	Sec./word	Words	Delete rate	Key/char	Paste	Cut	AI sim.	Typos	TTR	Yule's K	N-gram nov.	Peer orig.	Creativ. index
AI agent	-1.439***	-0.073	-1.551***	-1.890***	-0.334***	-0.151**	0.978***	-1.228***	0.828***	-1.048***	-0.392***	-1.167***	-1.130***
	(0.053)	(0.062)	(0.048)	(0.026)	(0.073)	(0.075)	(0.049)	(0.060)	(0.055)	(0.053)	(0.060)	(0.048)	(0.048)
Topic: Dreams	0.054	-0.044	-0.102	-0.061	0.037	0.133	0.201*	0.035	0.001	0.077	-0.270**	0.152	0.022
	(0.073)	(0.120)	(0.074)	(0.039)	(0.104)	(0.125)	(0.103)	(0.095)	(0.104)	(0.086)	(0.132)	(0.104)	(0.103)
Topic: Kindness	-0.012	0.150	-0.107	-0.075*	0.180	-0.004	0.230***	0.084	-0.358***	0.381***	0.242**	0.038	-0.064
	(0.068)	(0.120)	(0.079)	(0.044)	(0.133)	(0.009)	(0.086)	(0.097)	(0.101)	(0.093)	(0.105)	(0.095)	(0.090)
Topic: Music	-0.049	-0.133	0.008	-0.018	0.030	0.004	-0.103	0.116	0.215**	-0.199**	0.216*	-0.031	-0.231**
	(0.067)	(0.123)	(0.088)	(0.047)	(0.108)	(0.010)	(0.098)	(0.099)	(0.101)	(0.088)	(0.120)	(0.102)	(0.102)
Topic: Sleep	0.103	-0.272**	-0.061	-0.047	0.059	0.006	0.326***	0.035	0.180*	-0.052	0.097	0.086	-0.204*
	(0.097)	(0.112)	(0.081)	(0.044)	(0.107)	(0.009)	(0.121)	(0.094)	(0.103)	(0.087)	(0.120)	(0.095)	(0.105)
Topic: Smelling	-0.084	-0.139	-0.011	-0.004	-0.042	0.117	-0.194**	-0.028	0.023	0.091	0.515***	0.153	0.093
	(0.069)	(0.118)	(0.083)	(0.043)	(0.087)	(0.115)	(0.087)	(0.096)	(0.105)	(0.089)	(0.108)	(0.093)	(0.089)
Topic: Sugar	0.139	-0.010	-0.027	-0.030	0.088	0.211	-0.020	0.134	-0.253**	0.514***	0.379***	0.077	0.011
	(0.086)	(0.108)	(0.086)	(0.049)	(0.116)	(0.149)	(0.081)	(0.096)	(0.103)	(0.097)	(0.101)	(0.094)	(0.088)
Constant	0.794***	0.104	0.926***	1.109***	0.138	0.019	-0.624***	0.645***	-0.435***	0.469***	0.051	0.595***	0.694***
	(0.064)	(0.081)	(0.078)	(0.042)	(0.091)	(0.042)	(0.059)	(0.086)	(0.071)	(0.062)	(0.082)	(0.065)	(0.062)
N	955	955	955	955	955	955	955	955	955	955	955	955	955
Positive class N	543	543	543	543	543	543	543	543	543	543	543	543	543
Control N	412	412	412	412	412	412	412	412	412	412	412	412	412
R^2	0.514	0.018	0.593	0.878	0.033	0.012	0.269	0.375	0.215	0.333	0.097	0.336	0.327
Adj. R^2	0.510	0.010	0.590	0.877	0.025	0.005	0.263	0.370	0.209	0.328	0.090	0.331	0.322
Topic FE	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes

Notes: Standard errors in parentheses; HC1 robust SEs. All regressions include topic fixed effects, standardized outcomes, and omit the Coloring topic. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

Indicator	AI only				Any AI				AI assistance or Agent			
Indicator	ROC-AUC	Avg. prec.	Recall	F1	ROC-AUC	Avg. prec.	Recall	F1	ROC-AUC	Avg. prec.	Recall	F1
Behavior												
Seconds per word	0.601	0.642	0.668	0.591	0.586	0.831	0.669	0.712	0.836	0.939	0.779	0.816
Words	0.640	0.608	0.649	0.643	0.597	0.802	0.567	0.667	0.603	0.744	0.707	0.748
Deletion rate	0.762	0.806	0.697	0.681	0.730	0.898	0.683	0.754	0.883	0.944	0.853	0.874
Keystroke / char ratio	0.767	0.820	0.567	0.702	0.734	0.904	0.513	0.669	0.896	0.959	0.802	0.883
Paste detected	0.653	0.675	0.385	0.537	0.628	0.835	0.322	0.482	0.551	0.743	0.167	0.281
Cut detected	0.499	0.501	0.995	0.669	0.499	0.750	0.995	0.855	0.502	0.700	0.998	0.824
AI signature												
Similarity to AI baseline	0.724	0.793	0.565	0.646	0.668	0.879	0.520	0.648	0.737	0.905	0.643	0.759
Evaluation												
Typos flag	0.609	0.571	0.745	0.668	0.600	0.794	0.723	0.767	0.686	0.791	0.889	0.847
Creativity and originality												
Lexical TTR (higher = more diverse)	0.549	0.531	0.543	0.547	0.565	0.780	0.557	0.654	0.655	0.807	0.569	0.662
Lexical Yule's K (lower = more diverse)	0.637	0.628	0.668	0.633	0.610	0.821	0.640	0.714	0.740	0.867	0.712	0.768
N-gram novelty share (higher = more novel)	0.608	0.609	0.490	0.542	0.558	0.793	0.449	0.575	0.623	0.794	0.589	0.683
Peer originality (higher = more original)	0.699	0.778	0.567	0.622	0.633	0.865	0.521	0.639	0.747	0.906	0.644	0.750
Creativity index (higher = more creative)	0.723	0.795	0.575	0.642	0.656	0.875	0.517	0.643	0.748	0.907	0.642	0.758

Table A.3: Standalone predictive power of indicators across AI-use classifiers

Notes: Each cell reports performance from a standalone cross-validated logistic-regression classifier estimated with a single indicator at a time. Columns report ROC-AUC, average precision, recall, and F1. *AI only* compares human `AI_only` responses to Baseline Control; *Any AI* compares all human AI-assistance arms (`AI_only`, `AI_low`, and `AI_long_term`) to Baseline Control; *AI assistance or Agent* combines human `AI_only` responses and all AI agent responses against Baseline Control. Indicators are grouped into *Behavior*, *AI signature*, *Evaluation*, and *Creativity and originality*, matching the summary-only benchmark coefficient plot. Baseline Control has $N = 412$; positive-class sample sizes are $N = 416$ for AI only, $N = 1235$ for Any AI, and $N = 959$ for AI assistance or Agent.

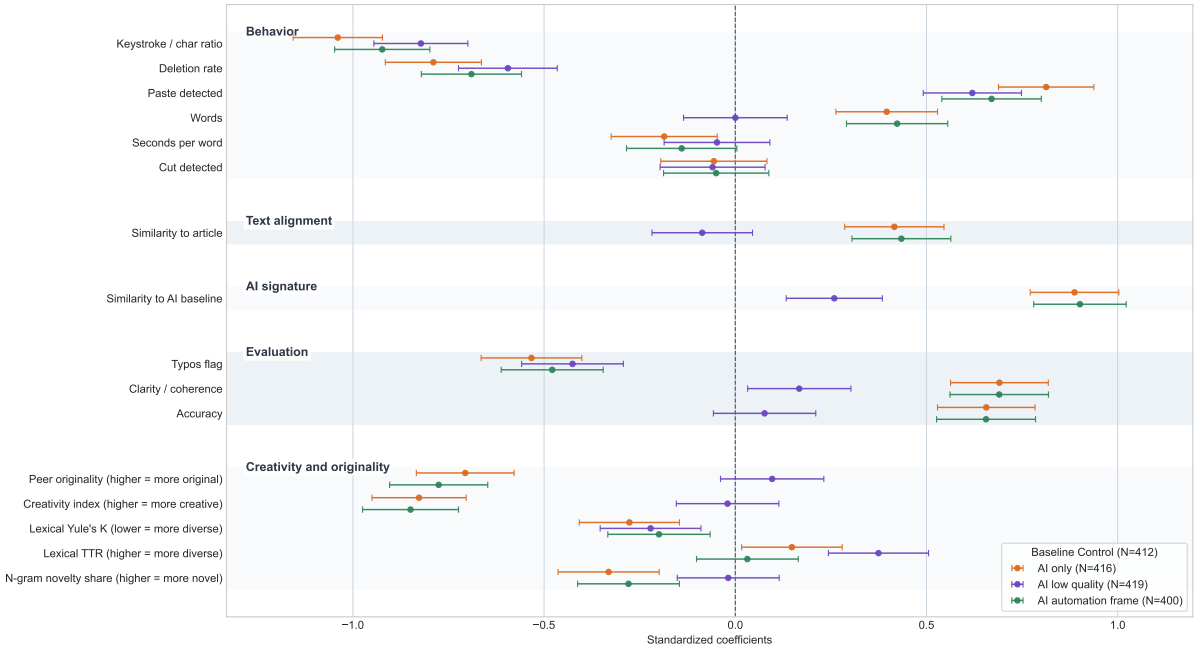

Figure A.4: Extended benchmark coefficients for human AI-treatment arms relative to baseline control.

Notes: The figure reports standardized coefficients from separate topic-adjusted OLS regressions estimated on the original human dataset, with the same baseline control group used as the reference category in each comparison. The three series correspond to `AI_only` ($n = 416$), `AI_low` ($n = 419$), and `AI_long_term` ($n = 400$), each compared against baseline control ($n = 412$). Outcomes are grouped into the same extended categories as in the main coefficient plot: *Behavior*, *Text alignment*, *AI signature*, *Evaluation*, and *Creativity and originality*. The extended specification includes *Seconds per word* in Behavior, *Similarity to article* in Text alignment, and *Clarity / coherence* and *Accuracy* in Evaluation. Whiskers denote 95% confidence intervals. Positive values indicate higher levels of the measure relative to baseline control.

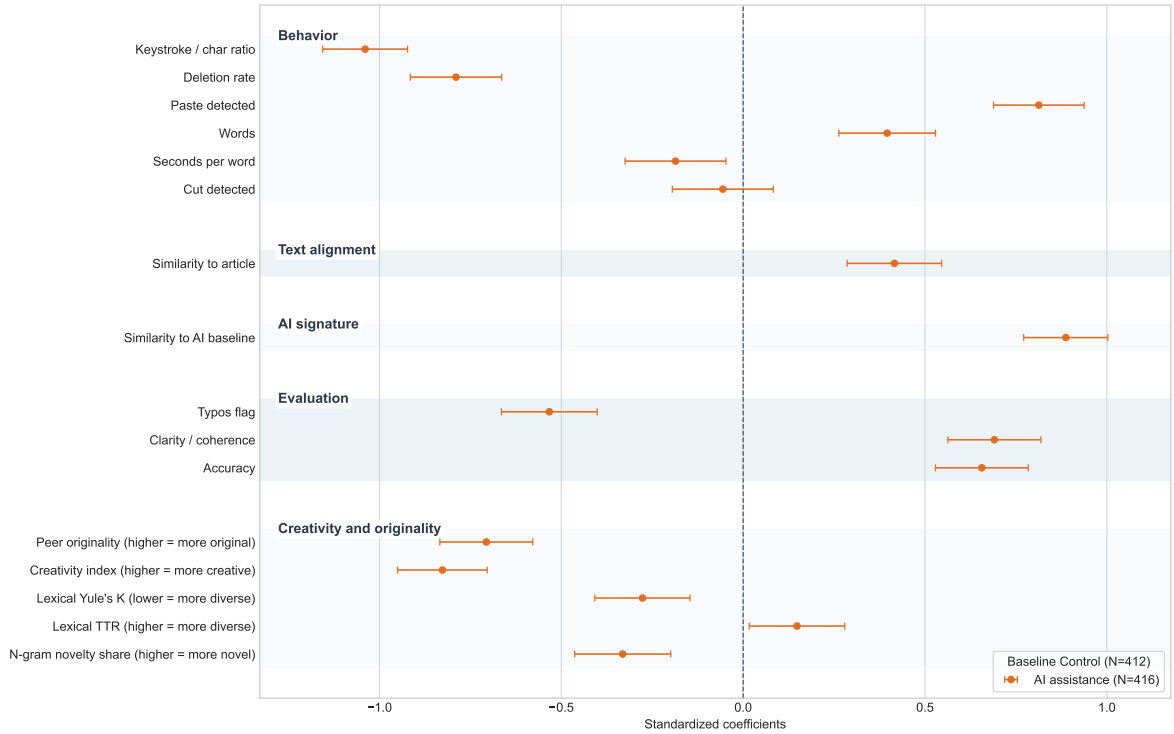

Figure A.5: Extended standardized differences in summary-task outcomes between the human AI-only treatment and the human control group.

Notes: This figure uses the same sample and specification as Figure 2: each coefficient comes from a separate topic-adjusted OLS regression with article-topic fixed effects, estimated on human `AI assisted` observations and control observations. Dependent variables are standardized, so coefficients are directly comparable across rows. Positive coefficients indicate that the AI-only group scores higher than the control group on the relevant outcome; negative coefficients indicate lower values. Relative to Figure 2, this extended specification adds a *Text alignment* category, where *Similarity to article* measures how closely the submitted summary tracks the source article, and expands the *Evaluation* category to include *Clarity / coherence* and *Accuracy*, where higher values indicate better evaluator assessments. As in Figure 2, the *Behavior* outcomes capture effort, speed, and copy/paste activity; the *AI signature* outcome captures similarity to an AI baseline; *Typos flag* is a binary indicator for whether typos were flagged; and the *Creativity and originality* outcomes capture originality, creativity, lexical diversity, and novelty.

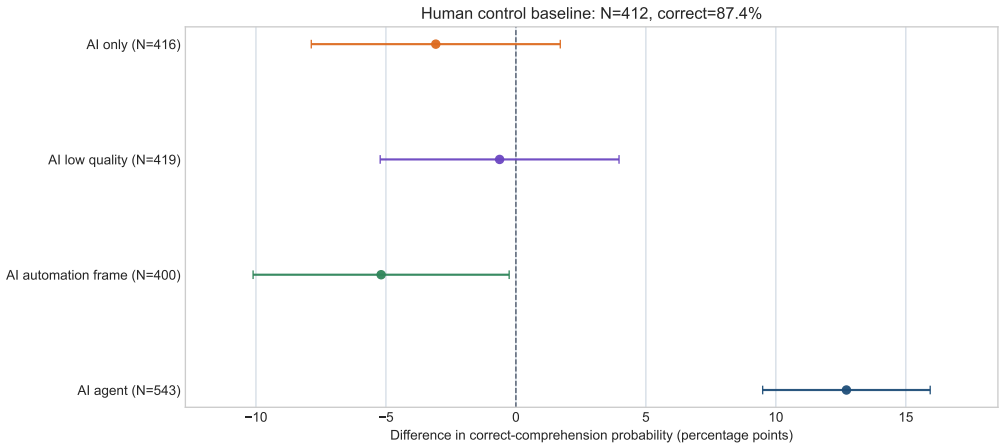

Figure A.6: Comprehension accuracy across AI conditions relative to human control.

Notes: Each point compares one AI condition with the shared human control baseline in a separate linear probability model with article-topic fixed effects. Horizontal bars show 95% HC1 confidence intervals. Coefficients are reported in percentage points, so positive values indicate higher probability of answering the comprehension question correctly than in the human control group and negative values indicate lower probability. The four contrasts are `AI_only`, `AI_low`, `AI_long_term`, and the AI-agent benchmark. The human control baseline contains $N = 412$ observations and an 87.4% correct-comprehension rate.

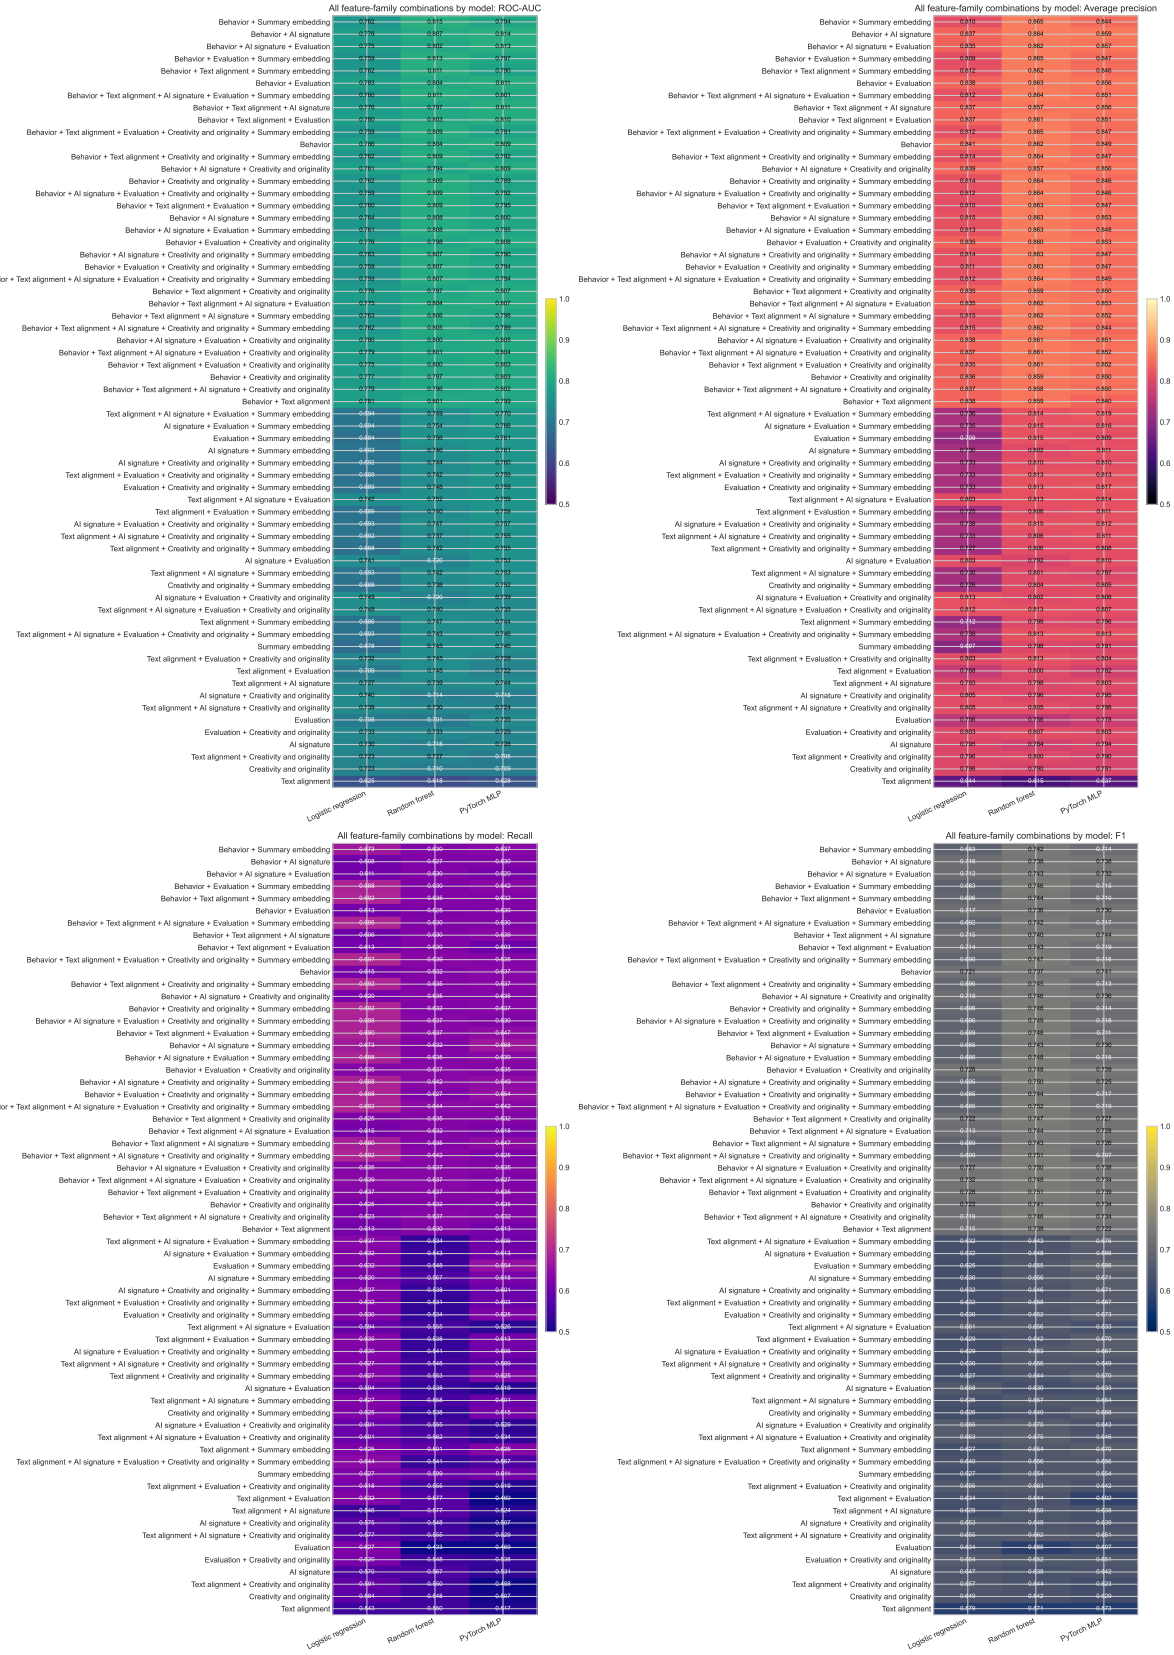

Figure A.7: AI-only assistance versus baseline control: classifier performance across feature-family combinations.

Notes: The figure reports out-of-sample classifier performance for distinguishing human responses written with AI assistance (`AI_only`) from baseline human control responses. The analysis sample contains $N = 828$ observations ($n = 416$ AI-only, $n = 412$ control). Columns correspond to logistic regression, random forest, and PyTorch MLP models. Rows correspond to all possible combinations of the feature families *Behavior*, *Text alignment*, *AI signature*, *Evaluation*, *Creativity and originality*, and *Summary embedding*. The four panels report ROC-AUC, average precision, recall, and F1. Rows are ordered by the highest ROC-AUC achieved by any model for that feature-family combination. The *Summary embedding* family is based on a 384-dimensional sentence-transformer representation of the written summary (`all-MiniLM-L6-v2`). In this comparison, the strongest-performing bundle combines *Behavior* and *Summary embedding*.

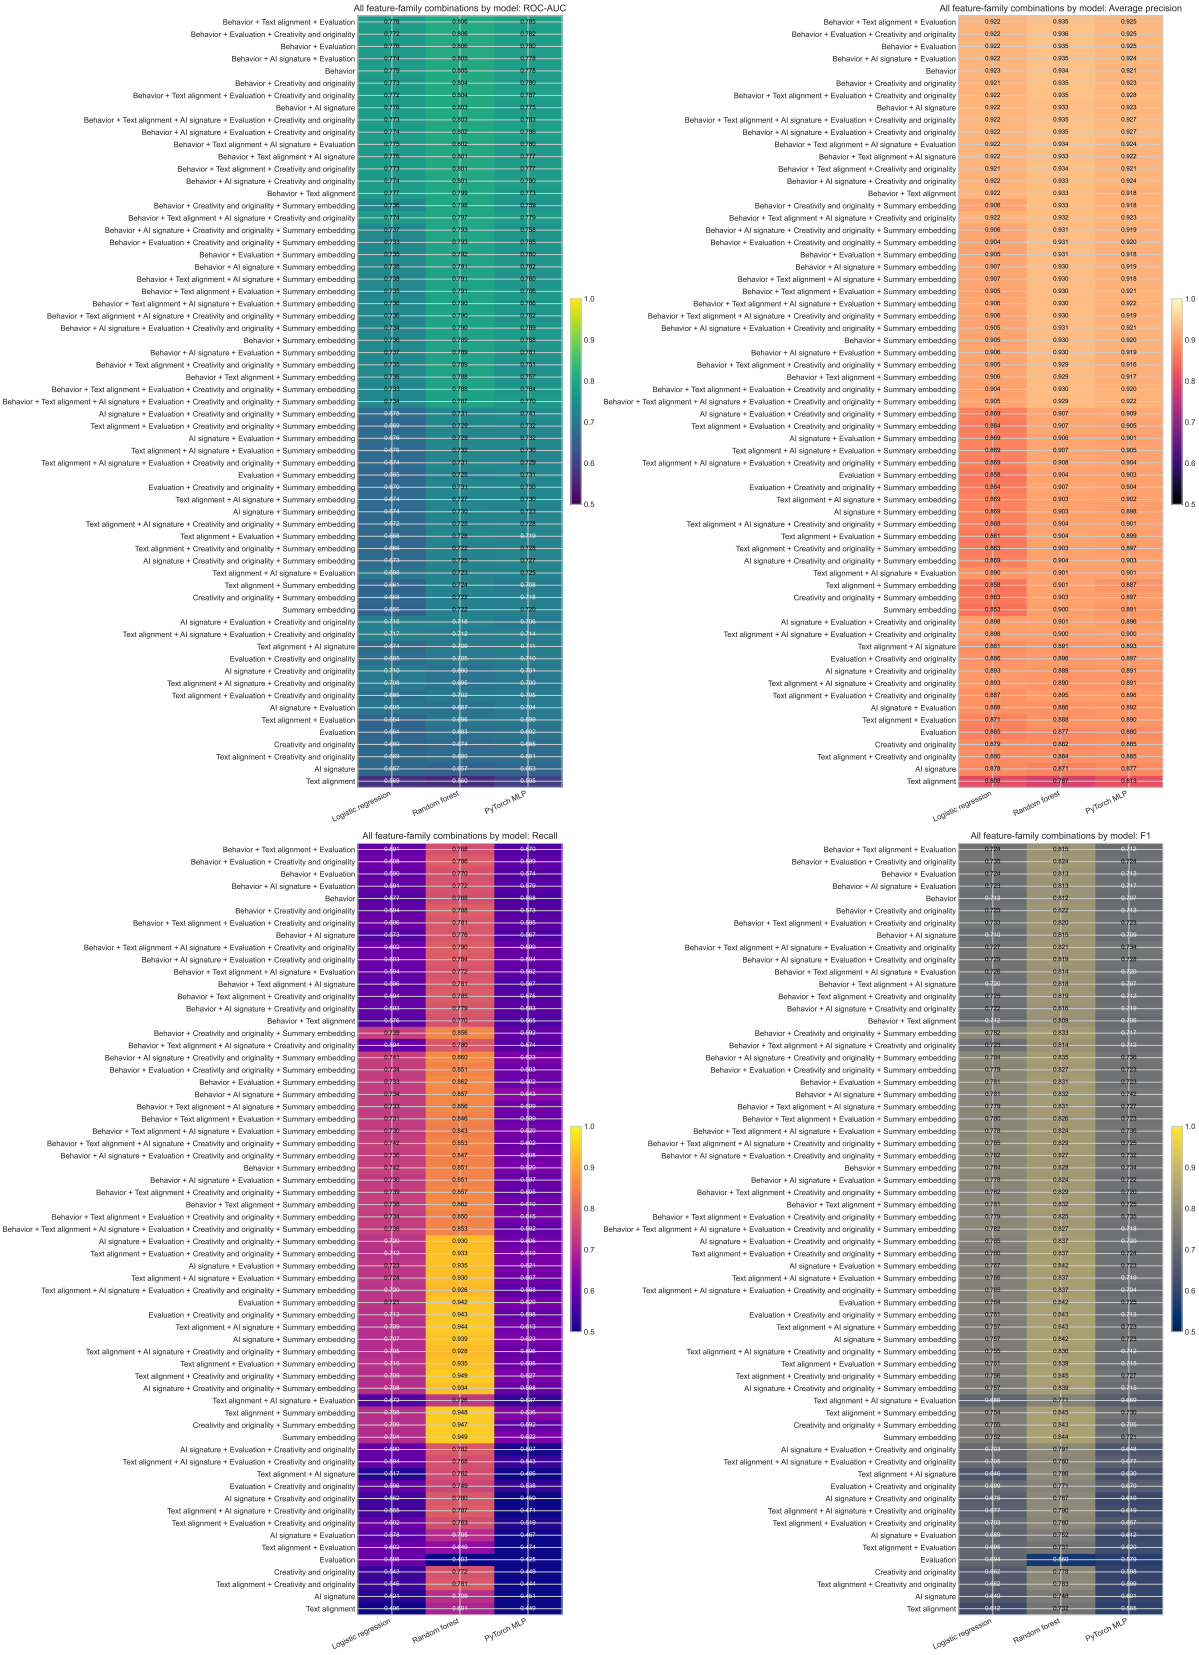

Figure A.8: Any AI assistance versus baseline control: classifier performance across feature-family combinations.

Notes: The figure reports out-of-sample classifier performance for distinguishing any AI-assisted human response (including those receiving low quality AI) from baseline human control responses. The analysis sample contains $N = 1647$ observations ($n = 1235$ Any-AI, $n = 412$ control). Columns correspond to logistic regression, random forest, and PyTorch MLP models. Rows correspond to all possible combinations of the feature families *Behavior*, *Text alignment*, *AI signature*, *Evaluation*, *Creativity and originality*, and *Summary embedding*. The four panels report ROC-AUC, average precision, recall, and F1. Rows are ordered by the highest ROC-AUC achieved by any model for that feature-family combination. The *Summary embedding* family is based on a 384-dimensional sentence-transformer representation of the written summary (`all-MiniLM-L6-v2`). In this broader comparison, the best-performing bundle combines *Behavior*, *Text alignment*, and *Evaluation*, while adding *Summary embedding* does not improve the top-performing specification.

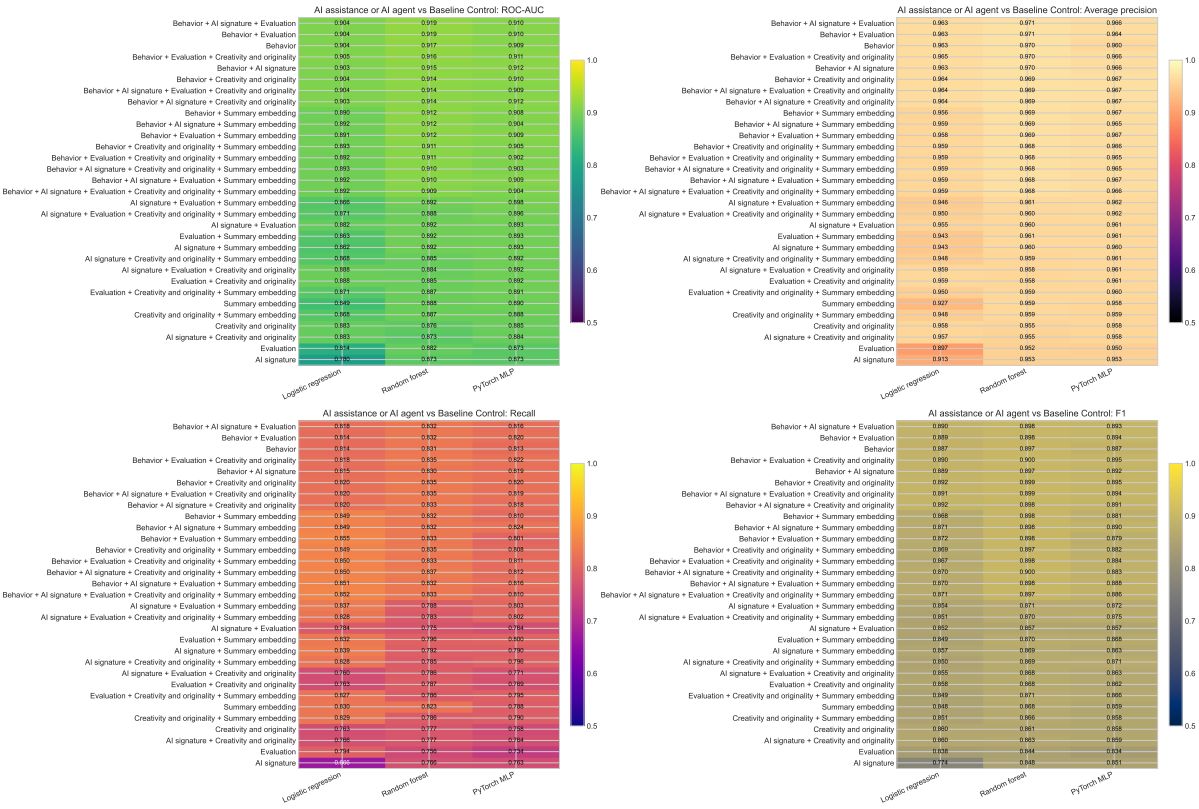

Figure A.9: Classifier performance across feature-family combinations: AI assistance or AI agent vs. Baseline Control

Notes: The positive class combines human `AI assistance` responses and all AI agent responses; the comparison group is Baseline Control. The four panels report cross-validated ROC-AUC, average precision, recall, and F1. Rows show feature-family combinations, and columns show the classifier used (logistic regression, random forest, and PyTorch MLP). Feature families are *Behavior*, *AI signature*, *Evaluation*, *Creativity and originality*, and *Summary embedding*. The summary embedding family is based on 384-dimensional summary-text embeddings from `sentence-transformers/all-MiniLM-L6-v2`. Larger values indicate better classifier performance. Sample sizes are $N = 412$ for Baseline Control and $N = 959$ for the combined AI group.

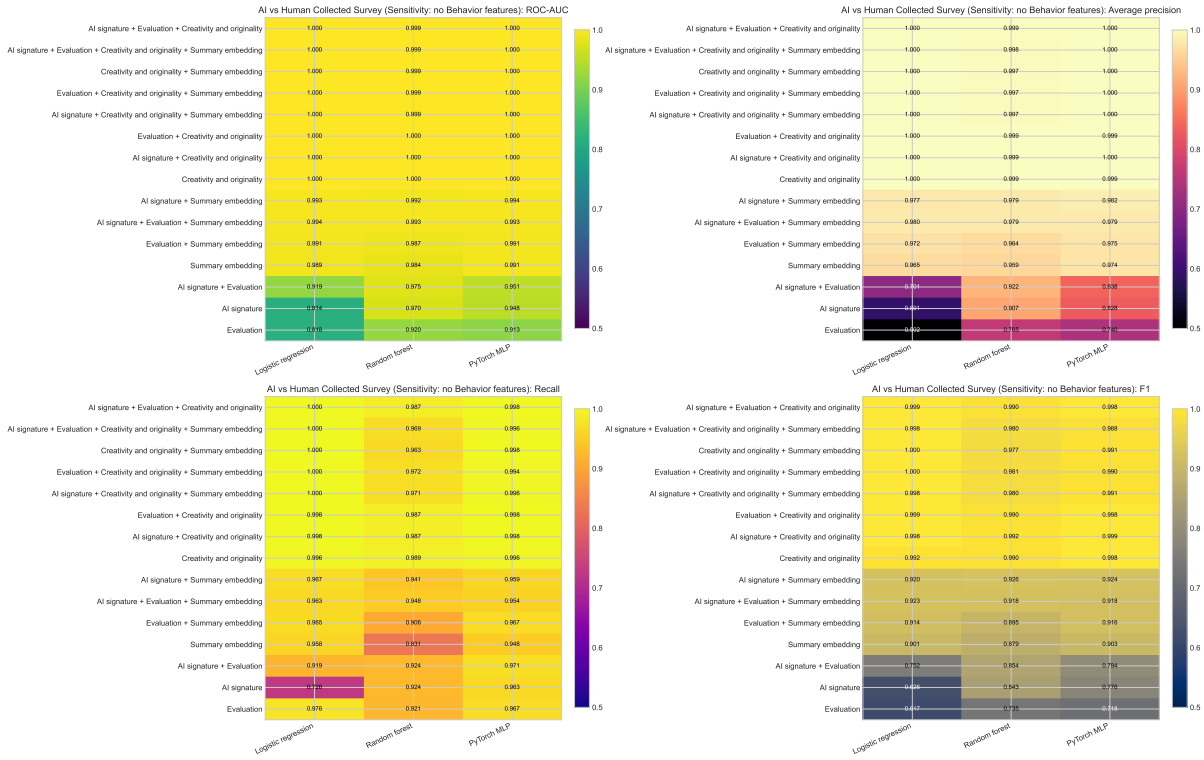

Figure A.10: AI agent versus human collected survey responses (sensitivity analysis excluding behavior features): classifier performance across feature-family combinations.

Notes: The figure reports out-of-sample classifier performance for distinguishing AI agent responses from the pooled human-collected survey sample after removing the full *Behavior* family. The human comparison group combines baseline control and all human AI-assistance arms. The analysis sample contains $N = 2190$ observations ($n = 543$ AI agent, $n = 1647$ human-collected survey). Columns correspond to logistic regression, random forest, and PyTorch MLP models. Rows correspond to all possible combinations of the remaining feature families: *AI signature*, *Evaluation*, *Creativity and originality*, and *Summary embedding*. The *Behavior* family is excluded entirely, so the models do not use timing, click, keystroke, or edit-telemetry variables. The four panels report ROC-AUC, average precision, recall, and F1. Rows are ordered by the highest ROC-AUC achieved by any model for that feature-family combination. The *Summary embedding* family is based on a 384-dimensional sentence-transformer representation of the written summary (`all-MiniLM-L6-v2`). In this sensitivity analysis, the first-ranked top-performing specification combines *Evaluation*, *Creativity and originality*, and *Summary embedding* with logistic regression, and multiple non-behavior feature bundles achieve near-perfect or perfect separation.