

The Political Ramification of AI-Augmented Work: Evidence from a Field Experiment *

Valentina González-Rostani[†] and Shir Raviv[‡]

May 8, 2026

Recent research on the political consequences of artificial intelligence has focused primarily on the threat of future job displacement. In this literature, workers are typically treated as passive recipients of technological change, responding to expert forecasts or firm-level adoption decisions. Yet this perspective misses a defining feature of the current AI wave. Unlike many earlier forms of automation, generative AI is already widely accessible at the individual level. Large language models are inexpensive, easy to use, and increasingly embedded in white-collar work. Workers can adopt them independently to summarize documents, draft text, generate ideas, and complete other cognitive tasks without employer direction. As a result, AI is not only a future threat but also a present workplace experience. This shift raises a different political question: how does direct engagement with AI-augmented work shape workers' behavior, perceptions of risk, and preferences over policy?

We argue that the political implications of AI depend not only on abstract expectations about long-run automation, but also on the concrete short-run experience of using AI tools on the job. These tools often produce immediate gains. They can reduce effort, save time, and improve performance. But those benefits may come with longer-term risks. Greater reliance on AI may weaken workers' confidence in their own skills, reduce autonomy, and gradually erode the comparative advantages that protect them from substitution. Moreover, workers who use these systems may contribute to their improvement and normalization, helping to develop technologies that could eventually replace human labor. Workers may therefore face a tension between immediate productivity gains and future vulnerability. Whether they perceive this tension, and whether they translate it into support for compensation, retraining, or restrictions on AI deployment, is an open empirical question.

Our framework identifies three mechanisms through which workers may learn about AI and update their political responses. First, the mere availability of AI assistance may serve as a signal that technological change is underway, making automation more salient even before workers use the tool. Second, direct experience with the quality of AI output may

*This study was pre-registered on 03/12/2026. The pre-analysis plan is available at: <https://aspredicted.org/i5jd26.pdf>.

[†]Department of Political Science, University of Southern California

[‡]Department of Political Science, Tel-Aviv University.

shape beliefs about technological capability. A high-performing system may make labor substitution more credible, while a low-performing system may reinforce the view that human judgment remains necessary. Third, workers may respond differently when the longer-term purpose of the technology is made explicit. Even if workers benefit from AI in the short run, they may not automatically connect those benefits to eventual labor replacement unless they are told that the technology is being developed for that purpose. These mechanisms suggest that political reactions to AI should depend on specific workplace experiences with the accessibility, performance, and framing of AI tools.

To examine these dynamics, we conduct a pre-registered field experiment on CloudResearch Connect, an online labor marketplace that enables the study of paid work under realistic task conditions. In March 2026, we recruited 1,675 U.S.-based workers to complete a writing assignment designed to resemble common forms of white-collar platform labor. Participants read a short factual article, write a summary of roughly 150 words, propose a concise title, and answer a comprehension question. These tasks require reading comprehension, prioritization, and clear communication, making them well suited to studying AI augmentation in routine cognitive work. To ensure that our results are not driven by the content of a particular text, participants are randomly assigned one of seven factual passages drawn from standardized exam materials and edited to be comparable in length and difficulty.

The experiment randomizes three features of workers' exposure to AI. First, participants are assigned either to a control condition with no AI access or to a treatment condition in which an AI writing assistant is embedded directly in the survey interface. Second, among those with AI access, we vary the quality of the tool. In the standard-performance condition, the assistant produces summaries that follow the task instructions and responds meaningfully to revision requests. In the low-performance condition, the same interface is paired with intentionally degraded output that is short, incomplete, stylistically inconsistent, and largely unresponsive to feedback. Third, among participants assigned to the standard-performance tool, we vary the framing of its purpose. Some workers are told only that the tool is being piloted to assist with the current task, while others are explicitly informed that the company plans eventually to automate these tasks and that workers' interactions with the tool will help improve the system toward that goal. This four-arm design allows us to distinguish the effect of AI access itself, the effect of learning from AI capability, and the effect of making long-run automation intentions salient.

We measure both workplace behavior and political attitudes. Our outcomes include support for compensatory labor-market policies such as financial assistance and retraining for workers displaced by new technologies; support for restrictive AI policies targeting frontier model development, firms' ability to automate tasks instead of using human workers, and individuals' use of AI applications; expectations about AI's effects on respondents, their communities, and the country over both short- and long-run horizons; and broader economic and social grievances related to status decline, corporate power, and future opportunity. Because the AI assistant is embedded in the study environment, we also observe actual patterns of use, including whether workers rely on the tool, how often they request revisions, and how they interact with it across rounds.

The study makes two broader contributions. First, it shifts research on the political economy of automation away from hypothetical scenarios and occupational exposure measures toward direct experience with AI in a realistic work setting. Second, it speaks to a

central puzzle in the literature: why perceived technological risk does not always generate stronger support for redistribution or regulation. Preliminary descriptive evidence already suggests that workers may distinguish between cushioning the costs of technological change and restricting the technology itself. More broadly, the paper argues that the political ramifications of AI will be shaped not only by forecasts of future displacement, but by workers' everyday experiences of using AI, learning from its capabilities and limits, and navigating the trade-off between immediate gains and longer-term vulnerability.